MySQL Basics & Commands

MySQL is a **relational database management system (RDBMS)** used to store, manage, and retrieve structured data efficiently. You interact with MySQL using SQL (Structured Query Language).

Logging into MySQL from Command Prompt (cmd)

Step 1: Open Command Prompt

- Windows: Press Win + R, type cmd, and hit Enter.
- Linux/macOS: Open the Terminal.

Step 2: Connect to MySQL

```
mysql -u root -p
```

- -u root → Specifies the username (root is the default admin user).
- -p → Prompts for a **password** (Enter the MySQL password when asked).

✓ Successful Login Output:

```
Welcome to the MySQL monitor. Commands end with ; or \g. mysql>
```

Now, you can execute MySQL commands.

2 MySQL Commands for Databases

1. Show all existing databases

SHOW DATABASES;

2. Create a new database

CREATE DATABASE mydatabase;

3. Use a specific database

Before working with tables, you need to **select** the database:

```
USE mydatabase;
```

Output:

Database changed

Now, all commands will apply to mydatabase.

4. Drop (Delete) a Database

DROP DATABASE mydatabase;

- This **permanently deletes** the database and all its tables.
- Be careful! This action cannot be undone.

Output:

```
Query OK, 1 row affected (0.01 sec)
```

4 Exiting MySQL

To exit MySQL and return to the command prompt, type:

EXIT;

QUIT;

Summary of Commands

Command

Description

mysql -u root -p Log in to MySQL

SHOW DATABASES; List all databases

CREATE DATABASE Create a new database

USE dbname; Switch to a database

DROP DATABASE Delete a database

dbname;

EXIT; Exit MySQL

Let me know if you need more details!

List of All Data Types in MySQL

MySQL has three main categories of data types:

- 1. Numeric Data Types
- 2. Date and Time Data Types
- 3. String (Character) Data Types

1. Numeric Data Types

Data Type

Data Type	Description
TINYINT	1-byte integer (-128 to 127)
SMALLINT	2-byte integer (-32,768 to 32,767)
MEDIUMINT	3-byte integer (-8,388,608 to 8,388,607)
INT/ INTEGER	4-byte integer (-2,147,483,648 to 2,147,483,647)
BIGINT	8-byte integer (-9 quintillion to 9 quintillion)
DECIMAL(m, d)	Exact fixed-point number (e.g., 10.2 means 10 digits, 2 decimal places)
FLOAT	Approximate floating-point number (4 bytes)

Description

DOUBLE / Approximate floating-point number (8 bytes)

REAL

BIT(n) Stores binary values (up to 64-bit)

2. Date and Time Data Types

Data Type Description

DATE Stores date (YYYY-MM-DD)

DATETIME Stores date and time (YYYY-MM-DD

HH:MM:SS)

TIMESTAM Stores timestamp (UTC-based,

p auto-updates)

TIME Stores time (HH:MM:SS)

YEAR Stores year (4 digits, e.g., 2025)

3. String (Character) Data Types

Data Type Description

CHAR(n) Fixed-length string (0-255 characters)

VARCHAR (Variable-length string (0-65,535 characters)

n)

TEXT Large text field (64 KB)

TINYTEXT Very small text (255 bytes)

MEDIUMTE Medium text (16 MB)

XT

LONGTEXT Large text (4 GB)

BLOB Binary Large Object (stores binary data)

TINYBLOB Small binary object (255 bytes)

MEDIUMBL Medium binary object (16 MB)

0B

```
LONGBLOB Large binary object (4 GB)

ENUM String with predefined values (ENUM('small', 'medium', 'large'))

SET Multiple predefined values (SET('a', 'b', 'c'))
```

1. Numeric Data Types Table

```
CREATE TABLE NumericData (
   id INT PRIMARY KEY AUTO_INCREMENT,
   tiny_int_col TINYINT,
   small_int_col SMALLINT,
   medium_int_col MEDIUMINT,
   int_col INT,
   big_int_col BIGINT,
   decimal_col DECIMAL(10,2),
   float_col FLOAT,
   double_col DOUBLE
);
```

Insert Data

```
INSERT INTO NumericData (
        tiny_int_col, small_int_col, medium_int_col, int_col, big_int_col,
        decimal_col, float_col, double_col
) VALUES
(127, 32000, 8000000, 2147483647, 9223372036854775807,
        12345.67, 123.45, 123456.789);
```

Check Table Structure

DESC NumericData;

2. Date and Time Data Types Table

```
CREATE TABLE DateTimeData (
```

```
id INT PRIMARY KEY AUTO_INCREMENT,
   date_col DATE,
   datetime_col DATETIME,
   timestamp_col TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
   time_col TIME,
   year_col YEAR
);

INSERT INTO DateTimeData (
   date_col, datetime_col, time_col, year_col
) VALUES
('2025-03-20', '2025-03-20 14:30:00', '14:30:00', 2025);
```

Check Table Structure

DESC DateTimeData;

3. String (Character) Data Types Table

```
CREATE TABLE StringData (
   id INT PRIMARY KEY AUTO_INCREMENT,
   char_col CHAR(10),
   varchar_col VARCHAR(50),
   text_col TEXT,
   tinytext_col TINYTEXT,
   mediumtext_col MEDIUMTEXT,
   longtext_col LONGTEXT,
   enum_col ENUM('low', 'medium', 'high'),
   set_col SET('A', 'B', 'C', 'D')
);

INSERT INTO StringData (
   char_col, varchar_col, text_col, tinytext_col, mediumtext_col,
longtext_col,
```

```
enum_col, set_col
) VALUES
('Hello', 'Hello World', 'This is a TEXT column', 'Tiny text data',
  'Medium text data', 'Long text data', 'high', 'A,B');
```

Check Table Structure

DESC StringData;

Summary

- ✓ We created three separate tables for different data types.
- ✓ We inserted sample data into each table.
- ✓ The DESC command helps verify the structure of each table.

This approach keeps data well-organized and easy to manage. Let me know if you need any modifications! $\sqrt[4]{}$

Using the DESC Command

The DESC command is used to describe the structure of a table.

```
sql
Copy code
DESC DataTypesDemo;
```

Importance of DESC Command

- Shows column names, data types, and NULL constraints.
- Displays default values (e.g., DEFAULT CURRENT_TIMESTAMP).
- Indicates primary keys, indexes, and auto-increment properties.
- Helps in debugging table structures before performing operations.

SQL Constraints: Definition & Usage

SQL constraints are rules applied to table columns to ensure data integrity, accuracy, and reliability. Constraints restrict the type of data that can be inserted into a table.

Types of SQL Constraints

SQL provides the following constraints:

- 1. NOT NULL
- 2. UNIQUE
- 3. PRIMARY KEY
- 4. FOREIGN KEY
- 5. CHECK
- 6. **DEFAULT**
- 7. **AUTO_INCREMENT** (MySQL-specific)
- 8. **INDEX** (Not exactly a constraint but helps optimize searches)

1. NOT NULL Constraint

- Ensures that a column cannot store NULL values.
- Use case: When a column must always have a value (e.g., email, username).

```
CREATE TABLE Users (
   id INT PRIMARY KEY AUTO_INCREMENT,
   name VARCHAR(50) NOT NULL,
   email VARCHAR(100) NOT NULL
);
```

Here, name and email must always have values (i.e., they cannot be NULL).

2. UNIQUE Constraint

- Ensures that all values in a column are different.
- Use case: To prevent duplicate entries (e.g., unique email, unique username).

Example

```
Sql
Copy code
CREATE TABLE Employees (
    emp_id INT PRIMARY KEY AUTO_INCREMENT,
    email VARCHAR(100) UNIQUE
);
```

Here, email must be unique, so no two employees can have the same email.

3. PRIMARY KEY Constraint

- Uniquely identifies each record in a table.
- A **PRIMARY KEY** is a combination of NOT NULL + UNIQUE.
- Use case: To uniquely identify records in a table.

Example

```
sql
Copy code
CREATE TABLE Products (
    product_id INT PRIMARY KEY,
    name VARCHAR(100) NOT NULL,
    price DECIMAL(10,2)
```

```
);
```

- Here, product_id is the PRIMARY KEY, meaning:
 - Each product_id must be unique.
 - Each product_id cannot be NULL.
- Composite Primary Key (multiple columns can be a primary key):

sql

Copy code

```
CREATE TABLE Orders (
    order_id INT,
    product_id INT,
    quantity INT,
    PRIMARY KEY (order_id, product_id)
);
```

Here, both order_id and product_id together form a PRIMARY KEY.

4. FOREIGN KEY Constraint

- Links two tables together (parent-child relationship).
- The foreign key column must refer to a PRIMARY KEY in another table.
- Use case: Ensuring referential integrity (e.g., orders must be linked to valid customers).

Example

sql

Copy code

```
CREATE TABLE Customers (
    customer_id INT PRIMARY KEY,
    name VARCHAR(100) NOT NULL
);

CREATE TABLE Orders (
    order_id INT PRIMARY KEY,
    customer_id INT,
    order_date DATE,
    FOREIGN KEY (customer_id) REFERENCES Customers(customer_id)
);
```

- Here, Orders.customer_id is a FOREIGN KEY, which means:
 - customer_id must exist in the Customers table.
 - If a customer is deleted, we need to **handle dependencies** (e.g., CASCADE, SET_NULL).
- Handling deletions:

sql

Copy code

FOREIGN KEY (customer_id) REFERENCES Customers(customer_id) ON DELETE CASCADE;

• If a customer is deleted, all their orders are deleted too.

sql

Copy code

```
FOREIGN KEY (customer_id) REFERENCES Customers(customer_id) ON DELETE SET NULL;
```

• If a customer is deleted, customer_id in Orders is set to NULL.

5. CHECK Constraint

- Ensures that column values meet a specific condition.
- Use case: To enforce business rules (e.g., salary should be > 0, age should be > 18).

Example

```
Sql
Copy code
CREATE TABLE Employees (
    emp_id INT PRIMARY KEY,
    name VARCHAR(50) NOT NULL,
    age INT CHECK (age >= 18),
    salary DECIMAL(10,2) CHECK (salary > 0)
);
```

- Here:
 - age must be 18 or more.
 - salary must be greater than 0.

6. DEFAULT Constraint

- Provides a default value if no value is given during insertion.
- Use case: To set default values (e.g., default status as "pending" for orders).

Example

```
copy code
CREATE TABLE Orders (
    order_id INT PRIMARY KEY,
    status VARCHAR(20) DEFAULT 'Pending'
);

• If we insert a new order without specifying status, it will automatically get the value 'Pending'.

sql
Copy code
INSERT INTO Orders (order_id) VALUES (1);
-- The status will be 'Pending'
```

7. AUTO_INCREMENT Constraint (MySQL only)

- Automatically generates unique numbers for each row.
- Use case: Used for primary keys to create unique IDs.

```
CREATE TABLE Employees (

emp_id INT PRIMARY KEY AUTO_INCREMENT,
```

```
name VARCHAR(100) NOT NULL
);

• Here, emp_id automatically increments for each new row.
sql
Copy code
INSERT INTO Employees (name) VALUES ('Alice');
```

The emp_id values will be 1 and 2 respectively.

INSERT INTO Employees (name) VALUES ('Bob');

8. INDEX (Not a constraint, but important)

- Speeds up searches and queries.
- **Use case:** When querying large tables frequently based on a specific column.

```
CREATE INDEX idx_name ON Employees(name);
```

Now, searching for employees by name will be faster.

Where and When to Use Constraints?

Constraint

Usage Scenario

NOT NULL

When a column **must always have a value** (e.g., username, email)

UNIQUE To ensure **no duplicate values** (e.g., email, phone number)

PRIMARY KEY To uniquely identify each row (e.g., user_id, order_id)

FOREIGN KEY To link tables (e.g., customer_id in Orders refers to

Customers)

CHECK To validate values (e.g., age >= 18, salary > 0)

DEFAULT To set default values (e.g., status as 'Pending')

AUTO_INCREMENT To generate unique IDs automatically (e.g., order_id,

emp_id)

INDEX To **speed up queries** (e.g., searching by name in Employees)

Foreign Key and Reference Key - Explanation & Usage

What is a Foreign Key?

A foreign key is a column (or set of columns) in one table that establishes a relationship with the primary key in another table. It ensures that the values in the foreign key column must exist in the referenced table.

What is a Reference Key?

A reference key is simply the primary key of the referenced table. The foreign key in the child table refers to this primary key in the parent table.

```
CREATE TABLE Departments (
```

```
dept_id INT PRIMARY KEY,
    dept_name VARCHAR(50)
);

CREATE TABLE Employees (
    emp_id INT PRIMARY KEY,
    emp_name VARCHAR(50),
    dept_id INT,
    FOREIGN KEY (dept_id) REFERENCES Departments(dept_id)
);
```

Here,

- dept_id in Departments is the reference key (primary key).
- dept_id in Employees is the foreign key that refers to the reference key.

Why is a Foreign Key Used?

- 1. Maintains Data Integrity
 - Ensures only valid data is inserted into the foreign key column.
 - Example: An employee must belong to an existing department.
- 2. Establishes Relationships Between Tables
 - Links tables logically, avoiding duplicate data.
 - Example: Employees belong to departments instead of repeating department names in every employee record.
- 3. Ensures Referential Integrity
 - Prevents orphan records (child records without a valid parent).
 - Example: You can't assign an employee to a non-existent department.

4. Helps in Cascading Operations

- Allows automatic updates or deletions when the parent record changes.
- Example: If a department is deleted, all employees in that department can also be deleted (ON DELETE CASCADE).

Summary:

Definition Term

Foreign Key Column in a child table that references a primary key in another table.

Reference The primary key in the parent table being referenced.

Key

Purpose Enforces data integrity, maintains relationships, and ensures

referential integrity.

Would you like more examples or a deeper explanation?

DDL (Data Definition Language) Commands in SQL

DDL (Data Definition Language) is used to define and manage the **structure** of database objects like tables, indexes, and constraints.



List of All DDL Commands in SQL

Command **Description** CREATE Creates **new database objects** (tables, views, indexes, etc.).

ALTER Modifies **existing database objects** (add, delete, modify columns).

DROP Deletes an **entire database object** (table, view, index).

TRUNCATE Deletes **all records** from a table but keeps the table structure.

RENAME Renames an existing database object.

COMMENT Adds a **comment** to database objects (table, column, etc.).

Explanation with Examples

1 CREATE – Create a Table

sql

);

```
Copy code

CREATE TABLE Employees (

emp_id INT PRIMARY KEY,

name VARCHAR(50),

age INT,

department VARCHAR(50)
```

• Creates a table Employees with emp_id, name, age, and department.

2 ALTER – Modify Table Structure

(a) Add a Column

sql

Copy code

ALTER TABLE Employees ADD salary DECIMAL(10,2);

Adds a new column salary to the Employees table.

(b) Modify Column Data Type

sql

Copy code

ALTER TABLE Employees MODIFY name VARCHAR(100);

Changes the data type of name to VARCHAR (100).

(c) Rename a Column

ALTER TABLE Employees CHANGE salary monthly_salary DECIMAL(10,2);

• Renames column salary to monthly_salary.

(d) Delete a Column

ALTER TABLE Employees DROP COLUMN age;

Removes the age column from the Employees table.

3 DROP – Delete an Object

DROP TABLE Employees;

• Deletes the entire table Employees permanently (all data and structure).

DROP DATABASE CompanyDB;

• Deletes the entire database CompanyDB.

4 TRUNCATE – Delete All Records

TRUNCATE TABLE Employees;

- Deletes all records from the Employees table but keeps the table structure.
- **★ Difference between DELETE and TRUNCATE:**

Feature	DELETE	TRUNCATE
Removes Specific Rows?	Yes (With WHERE)	X No (Deletes all rows)
Keeps Table Structure?	✓ Yes	✓ Yes
Can Be Rolled Back?	Yes (If within a transaction)	X No
Performance	Slower (logs each row deletion)	Faster (resets table directly)

5 RENAME – Rename Table or Column

RENAME TABLE Employees TO Staff;

• Renames table Employees to Staff.

6 COMMENT – Add Comments to Tables & Columns

COMMENT ON COLUMN Employees.name IS 'Employee Full Name';

Purpose

Adds a comment on the name column of Employees table.

Quick Recap

	Command	
1)1)1	Command	

CREATE Creates a new database object.

ALTER Modifies an **existing object** (add, modify, delete

column).

DROP Deletes an entire **object permanently**.

TRUNCATE Deletes all records but keeps the structure.

RENAME Renames a table or column.

Would you like me to cover DML (Data Manipulation Language) next? $\bigcirc \mathscr{A}$

UPDATE and DELETE Constraints on FOREIGN KEY in SQL

1. ON DELETE CASCADE

Effect:

Automatically deletes child records when the parent record is deleted.

Tables & Sample Data:

```
CREATE TABLE Departments (

dept_id INT PRIMARY KEY,

dept_name VARCHAR(50)
);

CREATE TABLE Employees (

emp_id INT PRIMARY KEY,

emp_name VARCHAR(50),

dept_id INT,

FOREIGN KEY (dept_id) REFERENCES Departments(dept_id) ON DELETE CASCADE
);

INSERT INTO Departments VALUES (1, 'HR'), (2, 'IT');

INSERT INTO Employees VALUES (101, 'Alice', 1), (102, 'Bob', 2), (103, 'Charlie', 1);
```

Operation:

DELETE FROM Departments WHERE dept_id = 1;

Output:

```
SELECT * FROM Employees;
-- | emp_id | emp_name | dept_id |
-- |------|
-- | 102 | Bob | 2 |
```

Explanation: Employees in dept_id = 1 (Alice and Charlie) are automatically removed.

2. ON DELETE SET NULL

Effect:

Sets the child foreign key to NULL when the parent record is deleted.

Tables & Sample Data:

```
CREATE TABLE Projects (

project_id INT PRIMARY KEY,

project_name VARCHAR(50)
);

CREATE TABLE Tasks (

task_id INT PRIMARY KEY,
```

```
task_name VARCHAR(50),

project_id INT,

FOREIGN KEY (project_id) REFERENCES Projects(project_id) ON DELETE SET NULL
);

INSERT INTO Projects VALUES (100, 'Website Redesign'), (200, 'App Launch');

INSERT INTO Tasks VALUES (1, 'Design Homepage', 100), (2, 'Test API', 200);
```

Operation:

DELETE FROM Projects WHERE project_id = 100;

Output:

```
SELECT * FROM Tasks;
-- | task_id | task_name | project_id |
-- |------|
-- | 1 | Design Homepage | NULL |
-- | 2 | Test API | 200 |
```

Explanation: The task linked to project_id = 100 keeps its data, but its project_id is set to NULL.

3. ON DELETE SET DEFAULT

Effect:

Resets the child foreign key to a default value when the parent record is deleted.

```
Tables & Sample Data:
```

```
CREATE TABLE Categories (
  category_id INT PRIMARY KEY DEFAULT 0,
  category_name VARCHAR(50)
);
CREATE TABLE Products (
  product_id INT PRIMARY KEY,
  product_name VARCHAR(50),
  category_id INT DEFAULT 0,
  FOREIGN KEY (category_id) REFERENCES Categories(category_id) ON DELETE SET
DEFAULT
);
INSERT INTO Categories VALUES (0, 'Uncategorized'), (1, 'Electronics');
INSERT INTO Products VALUES (500, 'Laptop', 1), (501, 'Desk', 0);
Operation:
DELETE FROM Categories WHERE category_id = 1;
Output:
SELECT * FROM Products;
-- | product_id | product_name | category_id |
-- |------|------|
```

```
-- | 500 | Laptop | 0 | -- | 501 | Desk | 0 |
```

Explanation: Products in the deleted category are reassigned to the default category_id = 0.

4. ON DELETE RESTRICT

Effect:

Prevents deletion if child records exist.

Tables & Sample Data:

```
CREATE TABLE Authors (
author_id INT PRIMARY KEY,
author_name VARCHAR(50)
);

CREATE TABLE Books (
book_id INT PRIMARY KEY,
book_title VARCHAR(50),
author_id INT,
FOREIGN KEY (author_id) REFERENCES Authors(author_id) ON DELETE RESTRICT
);

INSERT INTO Authors VALUES (1, 'J.K. Rowling');
INSERT INTO Books VALUES (100, 'Harry Potter', 1);
```

Operation:

DELETE FROM Authors WHERE author_id = 1;

Output:

-- X Error: "Cannot delete or update a parent row: a foreign key constraint fails"

Explanation: The deletion is blocked because the author has linked books.

Summary of FOREIGN KEY Constraints:

Constraint	Parent Operation	Child Effect
ON DELETE CASCADE	Delete	Child records are deleted
ON DELETE SET NULL	Delete	Child foreign key set to NULL
ON DELETE SET DEFAULT	Delete	Child foreign key set to default
ON DELETE RESTRICT	Delete	Parent deletion is blocked
ON DELETE NO ACTION	Delete	Same as RESTRICT (deletion blocked)
ON UPDATE CASCADE	Update	Child foreign key is updated

ON UPDATE SET NULL	Update	Child foreign key set to NULL
ON UPDATE SET DEFAULT	Update	Child foreign key set to default
ON UPDATE RESTRICT	Update	Parent update is blocked
ON UPDATE NO ACTION	Update	Same as RESTRICT (update blocked)

These examples illustrate how **foreign key constraints** help maintain data integrity in relational databases.

Here's an optimized version for **ON UPDATE** foreign key constraints, formatted for Google Docs:

Practical Examples of FOREIGN KEY Constraints – ON UPDATE

Foreign key constraints maintain data integrity when updating primary keys in relational databases. Below are examples with sample data and outputs for different **ON UPDATE** actions.

1. ON UPDATE CASCADE

Effect: Updates child foreign keys when the parent's primary key changes.

Tables & Data:

sql

Copy code

```
CREATE TABLE Students (

student_id INT PRIMARY KEY,

name VARCHAR(50)
```

```
);
CREATE TABLE Enrollments (
   enrollment_id INT PRIMARY KEY,
   student_id INT,
   course VARCHAR(50),
   FOREIGN KEY (student_id) REFERENCES Students(student_id) ON UPDATE
CASCADE
);
INSERT INTO Students VALUES (1, 'Alice');
INSERT INTO Enrollments VALUES (100, 1, 'Math');
Operation:
sql
Copy code
UPDATE Students SET student_id = 10 WHERE student_id = 1;
Output:
sql
Copy code
SELECT * FROM Enrollments;
-- | enrollment_id | student_id | course |
-- |-----|
```

Explanation: The student_id in the child table updates to match the new value in the parent table.

2. ON UPDATE SET NULL

Effect: Sets child foreign keys to NULL when the parent's primary key changes.

```
Tables & Data:
sql
Copy code
CREATE TABLE Teams (
    team_id INT PRIMARY KEY,
    team_name VARCHAR(50)
);
CREATE TABLE Members (
    member_id INT PRIMARY KEY,
    team_id INT,
    name VARCHAR(50),
    FOREIGN KEY (team_id) REFERENCES Teams(team_id) ON UPDATE SET NULL
);
INSERT INTO Teams VALUES (1, 'Developers');
INSERT INTO Members VALUES (101, 1, 'Bob');
```

```
Operation:
```

sql

Copy code

```
UPDATE Teams SET team_id = 100 WHERE team_id = 1;
```

Output:

sql

Copy code

Explanation: The team_id is set to NULL because the referenced value changed.

3. ON UPDATE SET DEFAULT

Effect: Sets child foreign keys to a default value when the parent's primary key changes.

Tables & Data:

sql

Copy code

```
CREATE TABLE Suppliers (
    supplier_id INT PRIMARY KEY DEFAULT 0,
    supplier_name VARCHAR(50)
```

```
);
CREATE TABLE Products (
    product_id INT PRIMARY KEY,
    supplier_id INT DEFAULT 0,
    FOREIGN KEY (supplier_id) REFERENCES Suppliers(supplier_id) ON
UPDATE SET DEFAULT
);
-- Insert default supplier
INSERT INTO Suppliers VALUES (0, 'No Supplier'), (1, 'Tech Corp');
INSERT INTO Products VALUES (500, 1), (501, 0);
Operation:
sql
Copy code
UPDATE Suppliers SET supplier_id = 100 WHERE supplier_id = 1;
Output:
sql
Copy code
SELECT * FROM Products;
-- | product_id | supplier_id |
-- |-----|
```

```
-- | 500 | 0 | -- Reset to default (No Supplier)
-- | 501 | 0 |
```

Explanation: The child records linked to the changed parent key are reset to the default value (0).

4. ON UPDATE RESTRICT

Effect: Prevents parent key updates if child records exist.

Tables & Data:

```
sql
Copy code
CREATE TABLE Users (
    user_id INT PRIMARY KEY,
    name VARCHAR(50)
);
CREATE TABLE Posts (
    post_id INT PRIMARY KEY,
    user_id INT,
    content TEXT,
    FOREIGN KEY (user_id) REFERENCES Users(user_id) ON UPDATE RESTRICT
);
INSERT INTO Users VALUES (1, 'Alice');
```

```
INSERT INTO Posts VALUES (100, 1, 'Hello World');

Operation:
sql
Copy code

UPDATE Users SET user_id = 10 WHERE user_id = 1;

Output:
sql
Copy code
-- X Error: "Cannot delete or update a parent row: a foreign key constraint fails"
```

Explanation: The update is blocked because the user has associated posts.

5. ON UPDATE NO ACTION

Effect: Same as RESTRICT – prevents updates if child records exist.

Tables & Data:

```
sql
Copy code
CREATE TABLE Employees (
    emp_id INT PRIMARY KEY,
    name VARCHAR(50)
);
```

```
CREATE TABLE Salaries (
    salary_id INT PRIMARY KEY,
    emp_id INT,
    amount DECIMAL(10,2),
    FOREIGN KEY (emp_id) REFERENCES Employees(emp_id) ON UPDATE NO
ACTION
);
INSERT INTO Employees VALUES (1, 'Alice');
INSERT INTO Salaries VALUES (100, 1, 5000);
Operation:
sql
Copy code
UPDATE Employees SET emp_id = 10 WHERE emp_id = 1;
Output:
sql
Copy code
-- X Error: "Cannot delete or update a parent row: a foreign key
constraint fails"
```

Explanation: The update is blocked because ON UPDATE NO ACTION prevents modification.

Summary of ON UPDATE Effects

Constraint	Parent Update Effect	Child Effect
ON UPDATE CASCADE	Parent key updated	Child records updated
ON UPDATE SET NULL	Parent key updated	Child foreign key set to NULL
ON UPDATE SET DEFAULT	Parent key updated	Child foreign key set to default
ON UPDATE RESTRICT	Parent key update blocked	No changes allowed
ON UPDATE NO ACTION	Parent key update blocked	No changes allowed

This structured breakdown makes it easy to understand how **ON UPDATE** constraints enforce referential integrity. Let me know if you need modifications!

Final Thoughts

- Understanding ON DELETE and ON UPDATE helps prevent orphan records and maintain consistency.
- Choosing the right constraint depends on business rules and data relationships.
- ✓ Use CASCADE carefully to avoid accidental mass deletions.

Let me know if you need further clarifications or code examples! 🚀

SQL INSERT, DELETE, and UPDATE with Multiple Examples

SQL provides three essential commands to manipulate data in tables:

- 1. INSERT → Adds new records.
- 2. UPDATE \rightarrow Modifies existing records.
- 3. DELETE → Removes records.

We will go through detailed examples with real-world scenarios.

1. INSERT Statement

Used to add new records into a table.

Syntax

sql

Copy code

```
INSERT INTO table_name (column1, column2, ...) VALUES (value1, value2,
...);
```

Example 1: Inserting a Single Record

```
sql
```

Copy code

```
CREATE TABLE Employees (
    emp_id INT PRIMARY KEY AUTO_INCREMENT,
    name VARCHAR(50) NOT NULL,
```

```
department VARCHAR(50),
    salary DECIMAL(10,2)
);

INSERT INTO Employees (name, department, salary)
VALUES ('Alice', 'HR', 50000.00);
```

emp_id is auto-generated, so we don't need to provide it.

Example 2: Inserting Multiple Records

```
INSERT INTO Employees (name, department, salary) VALUES
('John Doe', 'IT', 70000.00),
('Emma Watson', 'Finance', 60000.00),
('Mike Brown', 'IT', 75000.00);
```

✓ Inserting multiple rows in a single query improves performance.

Example 3: INSERT from Another Table

sql

Copy code

```
CREATE TABLE Employee_Backup AS
SELECT * FROM Employees WHERE department = 'IT';
```

This creates a backup table with only IT employees.

2. UPDATE Statement

Used to modify existing records.

Syntax

sql

Copy code

```
UPDATE table_name
SET column1 = value1, column2 = value2, ...
WHERE condition;
```



⚠ Without WHERE, all rows will be updated!

Example 1: Updating a Single Record

sql

Copy code

```
UPDATE Employees
SET salary = 80000
WHERE name = 'John Doe';
```

✓ John Doe's salary is now updated to 80,000.

Example 2: Updating Multiple Rows

sql

Copy code

```
UPDATE Employees

SET salary = salary * 1.10

WHERE department = 'IT';
```

✓ Increases the salary by 10% for all IT employees.

Example 3: Updating Based on Another Table

sql

Copy code

```
UPDATE Employees e

JOIN Employee_Backup b ON e.emp_id = b.emp_id

SET e.salary = b.salary * 1.05;
```

✓ Increases the salary of employees in the backup table by 5%.

3. DELETE Statement

Used to remove records.

Syntax

sql

Copy code

```
DELETE FROM table_name WHERE condition;
```



⚠ Without WHERE, all rows will be deleted!

Example 1: Deleting a Single Record

sql

Copy code

```
DELETE FROM Employees
WHERE name = 'Mike Brown';
```

Mike Brown is removed from the table.

Example 2: Deleting Multiple Rows

sql

Copy code

```
DELETE FROM Employees
WHERE department = 'HR';
```

All employees in HR are deleted.

Example 3: Deleting All Records (TRUNCATE)

sql

Copy code

TRUNCATE TABLE Employees;

Removes all records but keeps the table structure.

Example 4: DELETE with JOIN

sql

Copy code

DELETE e FROM Employees e
JOIN Employee_Backup b ON e.emp_id = b.emp_id;

✓ Deletes employees that exist in both Employees and Employee_Backup.

Final Thoughts

Command Use Case

INSERT Add new records

UPDATE Modify existing records

```
Remove specific records
 DELETE
 TRUNCATE Remove all records (faster than
              DELETE)

✓ Let me know if you need more practice queries or explanations! 

EXAMPLE:
use onetomany;
create table ss(
id int,
name varchar(10)
);
create table mm(
id int,
name varchar(10)
);
INSERT INTO ss (id, name) VALUES
(1, 'it'), (2, 'it'), (3, 'it'), (4, 'it'), (5, 'it'), (6, 'hr'), (7, 'sales'), (8, 'admin'), (9, 'finance'), (10,
'marketing');
insert into mm select *from ss where name='it';
select *from mm;
```

DELETE OPERATION DELETE ENTIRE ROW NOT SINGLE COLUMN USE UPDATE TO CHANGE OR DELETE ONE COLUMN

Here's the exact SQL script to create tables, insert data, and demonstrate the DELETE and UPDATE operations using three-table joins.

Create Tables and Insert Sample Data

```
sal
Copy code
-- Creating Customers Table
CREATE TABLE Customers (
    customer_id INT PRIMARY KEY,
    name VARCHAR(50),
    status VARCHAR(20) -- 'Active' or 'Inactive'
);
-- Creating Products Table
CREATE TABLE Products (
    product_id INT PRIMARY KEY,
    product_name VARCHAR(50),
    is_discontinued BOOLEAN -- 1 = discontinued, 0 = available
);
-- Creating Orders Table
CREATE TABLE Orders (
    order_id INT PRIMARY KEY,
    customer_id INT,
    product_id INT,
    order_date DATE
);
-- Insert Data into Customers
INSERT INTO Customers (customer_id, name, status) VALUES
(1, 'Alice', 'Active'),
(2, 'Bob', 'Inactive'),
(3, 'Charlie', 'Inactive');
-- Insert Data into Products
INSERT INTO Products (product_id, product_name, is_discontinued)
VALUES
```

```
(101, 'Laptop', 0),
(102, 'Keyboard', 1), -- Discontinued
(103, 'Mouse', 1); -- Discontinued

-- Insert Data into Orders
INSERT INTO Orders (order_id, customer_id, product_id, order_date)
VALUES
(1001, 1, 101, '2023-01-01'), -- Active customer, non-discontinued
product (retained)
(1002, 2, 102, '2023-02-01'), -- Inactive customer, discontinued
product (deleted)
(1003, 3, 103, '2023-03-01'); -- Inactive customer, discontinued
product (deleted)
```

2 DELETE Query with Three-Table JOIN

sql

```
Copy code
```

```
DELETE o FROM Orders o
JOIN Customers c ON o.customer_id = c.customer_id
JOIN Products p ON o.product_id = p.product_id
WHERE c.status = 'Inactive' -- Delete orders from inactive customers
   AND p.is_discontinued = 1; -- Delete orders for discontinued
products
```

Expected Result After Deletion:

sql

Copy code

SELECT * FROM Orders;

```
order_id customer_i product_id order_date
d

1001 1 101 2023-01-01
```



- Orders placed by inactive customers (Bob and Charlie) for discontinued products (Keyboard and Mouse) are deleted.
- Only the order for Alice (Active customer) with Laptop (not discontinued) remains.

3 Create More Tables for UPDATE Example

```
sql
Copy code
-- Creating Suppliers Table
CREATE TABLE Suppliers (
    supplier_id INT PRIMARY KEY,
    supplier_name VARCHAR(50),
    cost_rating VARCHAR(10) -- 'Low', 'Medium', 'High'
);
-- Creating Categories Table
CREATE TABLE Categories (
    category_id INT PRIMARY KEY,
    category_name VARCHAR(50)
);
-- Creating Products Table (Extended with supplier and category)
CREATE TABLE Products (
    product_id INT PRIMARY KEY,
    product_name VARCHAR(50),
    price DECIMAL(10,2),
    supplier_id INT,
    category_id INT
);
-- Insert Data into Suppliers
INSERT INTO Suppliers (supplier_id, supplier_name, cost_rating) VALUES
(1, 'TechGlobal', 'High'),
(2, 'EcoSupplies', 'Low');
-- Insert Data into Categories
INSERT INTO Categories (category_id, category_name) VALUES
(10, 'Electronics'),
```

```
(20, 'Accessories');
-- Insert Data into Products
INSERT INTO Products (product_id, product_name, price, supplier_id, category_id) VALUES
(101, '4K Monitor', 299.99, 1, 10), -- High-cost supplier, Electronics
(102, 'Wireless Mouse', 19.99, 2, 20), -- Low-cost supplier,
Accessories
(103, 'Gaming Chair', 199.99, 1, 10); -- High-cost supplier,
Electronics
```

4 UPDATE Query with Three-Table JOIN

```
sql
```

```
Copy code
UPDATE Products p
JOIN Suppliers s ON p.supplier_id = s.supplier_id
JOIN Categories c ON p.category_id = c.category_id
SET p.price = p.price * 0.9 -- Apply 10% discount
WHERE s.cost_rating = 'High' -- Target expensive suppliers
   AND c.category_name = 'Electronics'; -- Target Electronics category
```

▼ Expected Result After Update:

sql

Copy code

SELECT * FROM Products;

product_id	product_name	price	supplier_i d	category_i d
101	4K Monitor	269.99	1	10
102	Wireless Mouse	19.99	2	20
103	Gaming Chair	179.99	1	10

Explanation:

- Only products in the "Electronics" category and supplied by "High" cost suppliers (TechGlobal) are updated with a 10% price discount.
- The price of "4K Monitor" and "Gaming Chair" is reduced.
- The "Wireless Mouse" remains unchanged because its supplier has a low cost rating.

Summary of Key Concepts

Operation	Query	Purpose
DELETE	DELETE o FROM Orders o JOIN Customers c	Delete orders from inactive customers for discontinued products.
UPDATE	UPDATE Products p JOIN Suppliers s	Apply a 10% discount to expensive electronics from high-cost suppliers.

Now you can **run these queries in MySQL** and test them with the given sample data. Let me know if you need modifications or additional explanations! \mathscr{A}

The **SELECT** keyword in SQL is used to retrieve data from a database. Below is a comprehensive breakdown of its syntax, including clauses and optional components. The order of clauses **must follow the sequence outlined here**.

Basic Syntax

SELECT [DISTINCT | ALL] column1, column2, ...
FROM table_name
[WHERE condition]
[GROUP BY column(s)]
[HAVING condition]
[ORDER BY column(s) [ASC | DESC]]
[LIMIT n] -- MySQL/PostgreSQL
[OFFSET m] -- MySQL/PostgreSQL

SQL Operators Table

Category	Operator	Description	Example
----------	----------	-------------	---------

Comparison Operators	=	Equal to	WHERE salary = 50000
	i = or <>	Not equal to	WHERE department != 'IT'
	>	Greater than	WHERE salary > 60000
	<	Less than	WHERE hire_date < '2022-01-01'
	>=	Greater than or equal to	WHERE age >= 30
	<=	Less than or equal to	WHERE salary <= 45000
Logical Operators	AND	All conditions must be true	WHERE department = 'IT' AND salary > 60000
	OR	At least one condition is true	<pre>WHERE department = 'HR' OR department = 'Sales'</pre>
	NOT	Negates a condition	WHERE NOT department = 'IT'
Arithmetic Operators	+	Addition	SELECT salary + 5000 AS new_salary
	-	Subtraction	SELECT salary - tax AS net_salary
	*	Multiplication	SELECT quantity * price AS total
	/	Division	SELECT salary / 12 AS monthly_pay
Pattern Matching	%	Matches any sequence of characters	WHERE name LIKE 'A%' (Starts with "A")
	-	Matches a single character	WHERE name LIKE '_ob' (e.g., "Bob")
Membership Operators	IN	Value exists in a list	WHERE department IN ('IT', 'HR')
	NOT IN	Value does not exist in list	WHERE department NOT IN ('Sales')
Range Operators	BETWEEN	Between two values (inclusive)	WHERE salary BETWEEN 45000 AND 60000

	NOT BETWEEN	Outside a range	WHERE hire_date NOT BETWEEN '2020' AND '2022'
NULL Check Operators	IS NULL	Checks for NULL	WHERE hire_date IS NULL
	IS NOT NULL	Checks for non-NULL	WHERE department IS NOT NULL
String Operators			or+`
Set Operators	UNION	Combines results (removes duplicates)	SELECT name FROM IT UNION SELECT name FROM HR
	UNION ALL	Combines results (allows duplicates)	SELECT name FROM IT UNION ALL SELECT name FROM HR
	INTERSECT	Returns common rows	SELECT name FROM IT INTERSECT SELECT name FROM HR
	EXCEPT	Returns rows from the first query not in the second	SELECT name FROM IT EXCEPT SELECT name FROM HR
Bitwise Operators	&	Bitwise AND	SELECT 5 & 3 (Result: 1)
	•	•	Bitwise OR
	٨	Bitwise XOR	SELECT 5 ^ 3 (Result: 6)
Special Operators	EXISTS	Checks if a subquery returns rows	WHERE EXISTS (SELECT 1 FROM employees)
	DISTINCT	Returns unique values	SELECT DISTINCT department FROM employees
	CASE	Conditional logic	SELECT CASE WHEN salary > 50000 THEN 'High' ELSE 'Low' END

Pattern matching in SQL is primarily done using the LIKE operator, which is used to filter results based on a specific pattern. The two main wildcard characters used in LIKE are:

- % (percent sign) Matches **zero**, **one**, **or multiple** characters.
- _ (underscore) Matches a single character.

Sample Table: employees

We will use this table for examples:

id	name	department	email
1	Alice	IT	alice@company.com
2	Bob	HR	bob.hr@company.com
3	Charlie	IT	charlie@company.com
4	Diana	Finance	diana@finance.com
5	Eve	HR	eve123@company.co m

1. Using % for Pattern Matching

Example 1: Names that start with 'A'

```
SELECT * FROM employees WHERE name LIKE 'A%';
```

• Explanation:

- The % after 'A' means "any number of characters after A".
- This will return names like Alice, Alex, and Andrew.

Example 2: Emails that contain 'company'

```
SELECT * FROM employees WHERE email LIKE '%company%';
```

• Explanation:

- %company% means "contains the word 'company' anywhere in the string".
- This will return emails like **alice@company.com**, **bob.hr@company.com**, and **eve123@company.com**.

2. Using _ for Single Character Matching

Example 3: Names with 'o' as the second letter

```
SELECT * FROM employees WHERE name LIKE '_o%';
```

- _o% means "any name where the second letter is 'o".
- This will return names like Bob but not Alice or Charlie.

Example 4: Emails that start with exactly four characters before '@'

```
SELECT * FROM employees WHERE email LIKE '____@%';
```

- ____@% means "exactly four characters before the '@' symbol".
- This will return emails like "diana@finance.com", but not "alice@company.com" because "alice" has 5 characters.

3. Combining % and _ for Complex Patterns

Example 5: Emails ending with '.com' but having a digit before '@'

```
SELECT * FROM employees WHERE email LIKE '%[0-9]@%.com';
```

• Explanation:

- %[0-9]@%.com means:
 - \circ % \rightarrow Any number of characters before the digit.
 - \circ [0-9] \rightarrow A digit must appear before @.
 - \circ % \rightarrow Any characters after @ (domain name).

- \circ .com \rightarrow Ends with ".com".
- This will return **eve123@company.com**, but **not** "alice@company.com" (since Alice's email has no digit).

Output:

id	name	department	email
5	Eve	HR	eve123@company.co m

4. Using NOT LIKE for Exclusion

Example 6: Employees whose names do NOT start with 'A'

sql
CopyEdit
SELECT * FROM employees WHERE name NOT LIKE 'A%';

- Explanation:
 - NOT LIKE 'A%' excludes names that start with 'A'.

Conclusion

- % → Matches any number of characters (including zero).
- $\bullet \quad _ \to \mathsf{Matches} \; \textbf{a single character}.$
- $\bullet \quad \mathsf{LIKE} \to \mathsf{Finds} \ \mathsf{matches} \ \mathsf{based} \ \mathsf{on} \ \mathsf{a} \ \mathsf{pattern}.$
- NOT LIKE \rightarrow Excludes matches based on a pattern.

Let me know if you need more examples! \mathscr{A}

Membership Operators in SQL (IN and NOT IN) with SELECT

Membership operators are used to check if a value exists within a specified list of values. These operators include:

- IN → Checks if a value exists in a given list.
- NOT IN → Checks if a value does not exist in a given list.

Sample Table: employees

We will use this table for examples:

id	name	department	salary	hire_date	age
1	Alice	IT	50000	2021-03-15	28
2	Bob	HR	60000	2020-07-22	35
3	Charlie	IT	45000	2023-01-10	30
4	Diana	Finance	75000	2019-11-05	42
5	Eve	HR	48000	2022-09-30	29

1. Using IN Operator

Example 1: Selecting employees from IT and HR departments

```
SELECT * FROM employees WHERE department IN ('IT', 'HR');
```

• Explanation:

```
This query returns only those employees whose department is either IT or HR. SELECT * FROM employees WHERE department = 'IT' OR department = 'HR';
```



id	name	department	salary	hire_date	age
1	Alice	IT	50000	2021-03-15	28
2	Bob	HR	60000	2020-07-22	35
3	Charlie	IT	45000	2023-01-10	30
5	Eve	HR	48000	2022-09-30	29

Example 2: Selecting employees hired in specific years

SELECT * FROM employees WHERE YEAR(hire_date) IN (2020, 2021);

• Explanation:

• Filters employees hired in 2020 or 2021.

Output:

id	name	department	salary	hire_date	age
1	Alice	IT	50000	2021-03-15	28
2	Bob	HR	60000	2020-07-22	35

2. Using NOT IN Operator

Used to filter rows where the column value is **not** present in the given list.

Example 3: Selecting employees NOT from IT and HR departments

```
SELECT * FROM employees WHERE department NOT IN ('IT', 'HR');
```

Excludes employees from **IT** and **HR** departments.

Example 4: Selecting employees whose salary is NOT in a specific set

```
SELECT * FROM employees WHERE salary NOT IN (45000, 50000);
```

Filters out employees who have a salary of 45000 or 50000.

3. Using IN with Subqueries

We can also use IN with subqueries to match values dynamically.

Example 5: Select employees whose department has employees earning more than 60000

```
SELECT * FROM employees WHERE department IN
   (SELECT department FROM employees WHERE salary > 60000);
```

Explanation:

- Finds departments where at least one employee earns more than 60000.
- Returns all employees from those departments.

Output:

id	name	department	salary	hire_date	Ag e
2	Bob	HR	60000	2020-07-22	35
4	Diana	Finance	75000	2019-11-05	42
5	Eve	HR	48000	2022-09-30	29

Conclusion

- IN checks if a value exists in a given list.
- NOT IN checks if a value does not exist in a given list.
- IN can be used with **subqueries** to filter dynamically.

Would you like more complex queries? 🚀

Here are different **Range Operators (BETWEEN and NOT BETWEEN)** with explanations and examples using the SELECT statement.

1. Using BETWEEN with Numeric Values

Find employees with salaries between 45,000 and 60,000.

SELECT * FROM employees WHERE salary BETWEEN 45000 AND 60000;

- This query selects employees whose salary is between 45,000 and 60,000, inclusive of both values.
- Equivalent to:

```
SELECT * FROM employees WHERE salary >= 45000 AND salary
<=60000;</pre>
```

2. Using NOT BETWEEN with Numeric Values

Find employees whose salaries are NOT between 45,000 and 60,000.

```
SELECT * FROM employees
```

WHERE salary NOT BETWEEN 45000 AND 60000;

3. Using BETWEEN with Dates

₱ Find employees hired between 2020 and 2022.

```
SELECT * FROM employees
WHERE hire_date BETWEEN '2020-01-01' AND '2022-12-31';
```

4. Using NOT BETWEEN with Dates

Find employees hired BEFORE 2020 or AFTER 2022.

```
SELECT * FROM employees
WHERE hire_date NOT BETWEEN '2020-01-01' AND '2022-12-31';
```

5. Using BETWEEN with Age

Find employees aged between 28 and 35.

```
SELECT * FROM employees WHERE age BETWEEN 28 AND 35;
```

6. Using BETWEEN with String Values (Alphabetic Range)

*Find employees whose names fall alphabetically between 'B' and 'D'.

```
SELECT * FROM employees
WHERE name BETWEEN 'B' AND 'D';
```

7. Using NOT BETWEEN with String Values

Find employees whose names DO NOT fall alphabetically between 'B' and 'D'.

```
SELECT * FROM employees WHERE name NOT BETWEEN 'B' AND 'D';
```

Key Takeaways

- BETWEEN includes both **start** and **end** values in the range.
- NOT BETWEEN excludes the given range.
- Works with **numbers, dates, and strings**.
- ✓ Alphabetic comparison works in SQL but depends on collation settings.
- ✓ Use quotes (' ') for dates and string comparisons.

Would you like more advanced examples? **

Here are different **NULL Check Operators (IS NULL and IS NOT NULL)** with explanations and examples using the SELECT statement.

1. Using IS NULL

Find employees whose hire date is not available (NULL values).

SELECT * FROM employees WHERE hire_date IS NULL;

Explanation:

 This query selects employees where the hire_date column has NULL values (i.e., no data is recorded).

Output:

2. Using IS NOT NULL

Find employees who have a hire date (not NULL).

```
SELECT * FROM employees WHERE hire_date IS NOT NULL;
OUTPUT:
Id hire date
1 2-3-25
```

3. Using IS NULL with UPDATE

Assign a default hire date ('2024-01-01') to employees where hire_date is NULL.

```
UPDATE employees

SET hire_date = '2024-01-01'
WHERE hire_date IS NULL;
```

• Explanation:

• Updates all rows where hire_date is **NULL** and replaces it with '2024-01-01'.

4. Using IS NULL with DELETE

Remove all employees who don't have a hire date.

```
DELETE FROM employees
WHERE hire_date IS NULL;
```

Explanation:

 Deletes records where hire_date is NULL (i.e., employees without recorded hiring information).

5. Using IS NULL in an IF Condition

* Select employees and replace NULL hire dates with 'Not Available'.

• Explanation:

• The COALESCE() function replaces NULL values with 'Not Available'.

Output:

name	hire_date
Alice	2021-03-15
Bob	2020-07-22
Charli e	2023-01-10
Diana	2019-11-05
Eve	2022-09-30

Key Takeaways

- **IS NULL** checks for missing (NULL) values.
- IS NOT NULL filters out NULL values.
- ✓ Use COALESCE() to replace NULL with a default value.
- ✓ NULL is not the same as an empty string (' ') or zero (0).
- NULL values don't work with = or != (e.g., WHERE hire_date = NULL won't work!).

Set Operators in SQL with Examples (UNION, UNION ALL, INTERSECT, EXCEPT)

Set operators are used to **combine results** from two or more SELECT queries.

1. UNION (Removes Duplicates)

* Combine IT and HR employees into one list (without duplicates).

```
SELECT name, department FROM employees WHERE department = 'IT'
```

UNION

SELECT name, department FROM employees WHERE department = 'HR';

- Explanation:
 - Combines IT and HR department employees.
 - Removes duplicates automatically.

Output:

name departmen

t

Alice IT

Charli IT

е

Bob HR

Eve HR

2. UNION ALL (Keeps Duplicates)

★ Same as UNION, but keeps duplicates.

```
SELECT name, department FROM employees WHERE department = 'IT'
```

UNION ALL

```
SELECT name, department FROM employees WHERE department = 'HR';
```

Output (with duplicates if any exist):

```
name departmen
```

t

Alice IT

Charli IT

e

Bob HR

HR

Eve

3. INTERSECT (Common Records)

* Find employees working in both IT and HR departments.

```
SELECT name FROM employees WHERE department = 'IT'
INTERSECT
SELECT name FROM employees WHERE department = 'HR';
```

• Explanation:

- Returns only employees who exist in both queries.
- Some databases (like MySQL) don't support INTERSECT directly. You can simulate it with INNER JOIN.

Output:

name

(No common employees in this case)

4. EXCEPT (Records in First Query, But Not in Second)

* Find employees who are in IT but NOT in HR.

```
sql
CopyEdit
SELECT name FROM employees WHERE department = 'IT'
EXCEPT
SELECT name FROM employees WHERE department = 'HR';
```

- Explanation:
 - Returns employees in IT but NOT in HR.

Output:

name

Alice

Charlie

5. EXCEPT with UNION Example

Get employees from IT or HR, but exclude those from Finance.

```
SELECT name, department FROM employees WHERE department = 'IT'
UNION

SELECT name, department FROM employees WHERE department = 'HR'

EXCEPT

SELECT name, department FROM employees WHERE department = 'Finance';
```

Output:

```
namedepartmentAliceITCharlieITBobHREveHR
```

Key Takeaways

- UNION Combines results without duplicates.
- **UNION** ALL Combines results with duplicates.
- ✓ INTERSECT Returns only common records.
- **EXCEPT** Returns records from the first query that are NOT in the second.

Would you like advanced queries using set operators with filtering or sorting? 🚀

Step-by-Step Explanation of Set Operators in SQL

Set operators **combine results** from two or more SELECT statements. Here's how each command works internally:

1. UNION (Combines Results and Removes Duplicates

```
SELECT name, department FROM employees WHERE department = 'IT'
UNION
SELECT name, department FROM employees WHERE department = 'HR';
```

Step-by-Step Execution:

- First Query Execution (SELECT name, department FROM employees
 WHERE department = 'IT')
 - The database fetches all employees from the IT department.

```
Alice IT
Charlie IT
```

2.	Second Query Execution (SELECT	name,	department	FROM	employees
	WHERE department = 'HR')				

• The database fetches all employees from the HR department.

Bob HR

Eve HR

0

3. Combining Results (Before Removing Duplicates)

The database **merges** both result sets:

nginx

CopyEdit

Alice IT

Charlie IT

Bob HR

Eve HR

0

4. Removing Duplicates

 If there were duplicate values in both queries, UNION removes them automatically.

Final Output (After Deduplication)

nginx

CopyEdit

Alice IT

Charlie IT

Bob HR

Eve HR

5.

2. UNION ALL (Combines Results and Keeps Duplicates)

```
SELECT name, department FROM employees WHERE department = 'IT'
UNION ALL
SELECT name, department FROM employees WHERE department = 'HR';
```

Step-by-Step Execution:

- 1. **Executes the first SELECT query** to fetch IT employees.
- 2. Executes the second SELECT query to fetch HR employees.
- 3. Merges the results WITHOUT removing duplicates.

Final Output (Duplicates Remain):

```
Alice IT
```

Charlie IT

Bob HR

Eve HR

4.

Difference from UNION?

- UNION ALL is faster because it skips duplicate removal.
- Duplicates appear if there are any.

3. INTERSECT (Finds Common Records in Both Queries)

• Some databases (like MySQL) do not support INTERSECT directly, but we can simulate it using INNER JOIN.

```
SELECT name FROM employees WHERE department = 'IT'
INTERSECT
SELECT name FROM employees WHERE department = 'HR';
```

Step-by-Step Execution:

First Query Execution (SELECT name FROM employees WHERE department = 'IT')

Alice

Charlie

0

 Second Query Execution (SELECT name FROM employees WHERE department = 'HR')

Fetches HR employees:

Bob

Eve

0

3. Finding Common Records

- The database **compares both result sets** and finds names that exist in both.
- In this case, no common names exist.

Final Output:

```
(No results)

4.

If "Alice" worked in both IT and HR, the output would be:

nginx

CopyEdit

Alice
```

4. EXCEPT (Finds Records in First Query but NOT in Second Query)

Query:

```
sql
CopyEdit
SELECT name FROM employees WHERE department = 'IT'
EXCEPT
SELECT name FROM employees WHERE department = 'HR';
```

Step-by-Step Execution:

Executes the first query (SELECT name FROM employees WHERE department = 'IT')
Alice
Charlie
1.
<pre>Executes the second query (SELECT name FROM employees WHERE department = 'HR')</pre>
Bob
Eve
2.3. Removing Matching Records
 The database removes any records from the first result set that also appear in the second result set. Since "Alice" and "Charlie" do NOT appear in the second list, they remain.
Final Output:
Alice
Charlie
4.
If Alice were in HR too, she would be removed from the output.

5. EXCEPT with UNION Example

★ Get employees from IT or HR, but exclude Finance employees.

```
SELECT name, department FROM employees WHERE department = 'IT'
UNION

SELECT name, department FROM employees WHERE department = 'HR'

EXCEPT

SELECT name, department FROM employees WHERE department = 'Finance';
```

Step-by-Step Execution:

UNION (Combines IT and HR employees without duplicates)

Alice IT

Charlie IT

Bob HR

Eve HR

Diana Finance

1.

EXCEPT (Removes Finance employees from the above list)

Alice IT

Charlie IT

Bob HR

Eve HR

Key Takeaways

- ✓ UNION Combines results without duplicates.
- **UNION** ALL Combines results with duplicates.
- ✓ INTERSECT Returns only common records.
- **EXCEPT** Returns records from the first query that are NOT in the second.
- UNION ALL is faster than UNION because it skips duplicate checking.
- INTERSECT and EXCEPT work only in databases that support them (not MySQL).

Would you like examples using real-world scenarios like sales, customers, or orders? **

How EXCEPT Works with UNION in SQL?

Let's analyze how EXCEPT behaves when used with UNION. Specifically, does EXCEPT check both **name and department** or just one of them?

SQL Query:

sql

CopyEdit

SELECT name, department FROM employees WHERE department = 'IT'

UNION

SELECT name, department FROM employees WHERE department = 'HR'

EXCEPT

```
SELECT name, department FROM employees WHERE department =
'Finance';
Step-by-Step Execution:
Step 1: Execute First Query (IT Department)
SELECT name, department FROM employees WHERE department = 'IT';
Result Set 1 (IT employees):
Alice
          IT
Charlie
         IT
  1.
Step 2: Execute Second Query (HR Department)
sql
CopyEdit
SELECT name, department FROM employees WHERE department = 'HR';
Result Set 2 (HR employees):
nginx
CopyEdit
Bob
          HR
Eve
          HR
  2.
Step 3: Apply UNION (Combine IT and HR without Duplicates)
```

nginx

```
CopyEdit
Alice
          IT
Charlie IT
Bob
          HR
Eve
          HR
   3.
Step 4: Execute EXCEPT Query (Finance Employees)
SELECT name, department FROM employees WHERE department =
'Finance';
Finance Employees (To Be Excluded):
nginx
CopyEdit
Diana
          Finance
   4.
   5. Step 5: Apply EXCEPT
        • EXCEPT removes rows that are an exact match in both columns (name
           and department).
        o Since "Diana Finance" is not in the previous result set, no records are
           removed.
Final Output
nginx
CopyEdit
```

Alice

IT

Charlie IT
Bob HR
Eve HR

Does EXCEPT Check Both Name and Department?

Yes! EXCEPT compares both columns (i.e., name and department) together.

- If only the name matches but the department does not, the record is not removed.
- If both name and department match, the record is removed.

Example to Prove This:

If our Finance department had an employee also named **Charlie**, like this:

Charlie Finance

And our original dataset had:

Charlie IT

Then "Charlie Finance" would not remove "Charlie IT", because EXCEPT checks both columns together.

Key Takeaways

- **EXCEPT checks both columns** (not just one).
- ✓ If both name and department match, the record is removed.
- If only name matches but the department is different, the record remains.

Would you like me to demonstrate this with a modified dataset where EXCEPT actually removes some records?

Bitwise Operators in SQL

Bitwise operators perform operations at the **binary level** on integers. These are useful for tasks like permissions, flags, or low-level data manipulation.

1. List of Bitwise Operators in SQL

Operator	Description	Example	Result (Binary Calculation)
& (Bitwise AND)	Compares bits and returns 1 if both bits are 1	SELECT 5 & 3;	1 (0101 & 0011 = 0001)
•	` (Bitwise OR)	Compares bits and returns 1 if at least one bit is 1	`SELECT 5
^ (Bitwise XOR)	Returns 1 if bits are different (0 if same)	SELECT 5 ^ 3;	6 (0101 ^ 0011 = 0110)

~ (Bitwise NOT)	Flips all bits (1 becomes 0, 0 becomes 1)	SELECT ~5;	-6 (depends on system representation)
<< (Left Shift)	Shifts bits left by n places (multiplies by 2^n)	SELECT 5 << 1;	10 (0101 << 1 = 1010)
>> (Right Shift)	Shifts bits right by n places (divides by 2^n)	SELECT 5 >> 1;	2 (0101 >> 1 = 0010)

Operator	Description	Example	Output (5, 3)
&	Bitwise AND	SELECT 5 & 3;	1
		Bitwise OR	`SELECT 5
^	Bitwise XOR	SELECT 5 ^ 3;	6
~	Bitwise NOT	SELECT ~5;	-6
<<	Left Shift	SELECT 5 << 1;	10

|--|

Would you like more real-world use cases? 🚀

You said:

now do with special opertopr

ChatGPT said:

Special Operators in SQL

Special operators in SQL perform advanced filtering, existence checks, and conditional logic.

1. List of Special Operators

Operator	Description	Example
EXISTS	Checks if a subquery returns any rows	WHERE EXISTS (SELECT 1 FROM employees WHERE salary > 50000);
DISTIN CT	Returns unique values in a column	SELECT DISTINCT department FROM employees;
CASE	Implements conditional logic	SELECT name, salary, CASE WHEN salary > 50000 THEN 'High' ELSE 'Low' END AS salary_level FROM employees;

2. Examples with Explanation

Example 1: EXISTS (Checking for Existence in Another Table)

* Find employees who have made at least one sale.

sql

CopyEdit

SELECT name

FROM employees e

WHERE EXISTS (SELECT 1 FROM sales s WHERE s.employee_id = e.id);

How it works:

- The **subquery** (SELECT 1 FROM sales WHERE s.employee_id = e.id) checks if an employee has a sale.
- EXISTS returns **TRUE** if at least one row exists in sales.
- If EXISTS is **TRUE**, the employee is included in the result.

Output:

nginx

CopyEdit

name

Alice

Bob

Eve

(Only employees with sales are listed.)

Example 2: DISTINCT (Getting Unique Values)

PGet a list of unique departments from employees.

SELECT DISTINCT department FROM employees;

How it works:

Removes duplicate department values.

Example 3: CASE (Conditional Logic)

Classify employees based on salary.

```
SELECT name, salary,

CASE

WHEN salary > 60000 THEN 'High Salary'

WHEN salary BETWEEN 40000 AND 60000 THEN 'Medium Salary'

ELSE 'Low Salary'

END AS salary_category

FROM employees;
```

What is NOT EXISTS?

NOT EXISTS returns TRUE if the **subquery returns no rows**. It's typically used to **find records in one table that don't have matching entries in another**.

```
Syntax:
```

```
sql
CopyEdit
SELECT * FROM table1
WHERE NOT EXISTS (
    SELECT 1 FROM table2 WHERE table1.id = table2.id
);
```

⊚ Example Scenario:

Let's say we have two tables:

employees

```
emp_id name

1 Alice

2 Bob

3 Charli
e
```

attendances

```
emp_id date
```

```
1 2025-04-2
0 3 2025-04-2
```

? Goal:

Find employees who did NOT mark attendance on 2025-04-20

✓ Query Using NOT EXISTS:

```
sql
CopyEdit
SELECT * FROM employees e
WHERE NOT EXISTS (
    SELECT 1 FROM attendances a
    WHERE a.emp_id = e.emp_id
    AND a.date = '2025-04-20'
);
```

Output:

emp_id name

2 Bob

@ Explanation:

- It checks for each employee whether there's a matching entry in the attendances table for that date.
- If not found \rightarrow the employee is returned.

How it works:

- CASE checks each salary and assigns a category.
- If salary > 60000, output = "High Salary".
- If salary BETWEEN 40000 AND 60000, output = "Medium Salary".
- Otherwise, output = "Low Salary".

Output:

name	salary	salary_category
Alice	50000	Medium Salary
Bob	70000	High Salary
Charlie	40000	Medium Salary
Diana	80000	High Salary
Eve	35000	Low Salary

(Each employee is categorized.)

3. Summary of Special Operators

Operator	Purpose	Example Usage
EXISTS	Checks for existence of related data	<pre>WHERE EXISTS (SELECT 1 FROM sales WHERE employee_id = e.id);</pre>
DISTINCT	Removes duplicate values	SELECT DISTINCT department FROM employees;
CASE	Implements conditional logic	SELECT name, CASE WHEN salary > 50000 THEN 'High' ELSE 'Low' END FROM employees;

Use Case With a Table:

Let's say you have a students table:

SELECT name, marks, age FROM students WHERE marks > 60 AND age < 20;

This will fetch students with marks over 60 AND under 20 years of age.

JOINS

Table: Employees

EmpID	EmpName	DeptID
101	Alice	1
102	Bob	2
103	Charlie	3
104	David	NULL

Table: Departments

DeptID	DeptName
1	HR
2	IT
4	Marketing

Absolutely! Let's break down what's *actually happening under the hood* for each join type using the Employees and Departments tables.





- This is like saying:
 - "Combine every row from Employees with **every single** row from Departments, no conditions at all."
- It's like a cartesian product in math.

How it works:

- 4 Employees × 3 Departments = 12 combinations
- Doesn't care about DeptID, just blindly pairs each row.

2. THETA JOIN (using ON condition)

Process:

- This is a **conditional join** where you specify the logic for matching rows.
- In our case:
 "Only join rows where Employees.DeptID = Departments.DeptID."

How it works:

- It goes row-by-row and checks:

 "Hey, does this employee's DeptID match any department's DeptID?"
- If yes

 it joins them into a result row.
- If not 👉 it skips that employee.

Result:

• Only Alice (DeptID = 1) and Bob (DeptID = 2) have matching departments.

• Charlie (DeptID = 3) and David (NULL) are left out because there's no match.

Nocess:

- It's like a smart theta join. It automatically finds columns with the **same name** in both tables.
- Then it **implicitly** joins using those columns.

Under the hood:

SQL sees that both tables have a DeptID, so it assumes you want:

sql
CopyEdit
ON Employees.DeptID = Departments.DeptID

A Caution:

- If tables share multiple columns with the same name (e.g. both had Location), it will join on all of them.
- So while natural join is convenient, it can be **risky** if you're not careful.

99 Summary Table

Join Type Matches? Condition Rows Included Required?

Cross Join	No matching needed	X No	All combinations (cartesian)
Theta Join	Custom condition (=)	✓ Yes	Only matched rows based on logic
Natural Join	Auto match by column name	X No (implicit)	Only matched rows

Want to go next into **LEFT / RIGHT / FULL OUTER JOINs**? Those show what happens to unmatched data.



SELECT * FROM Employees CROSS JOIN Departments;

Output:

EmpID	EmpNam e	DeptID (Emp)	DeptID (Dept)	DeptNam e
101	Alice	1	1	HR
101	Alice	1	2	IT

101	Alice	1	4	Marketing
102	Bob	2	1	HR
102	Bob	2	2	IT
102	Bob	2	4	Marketing
103	Charlie	3	1	HR
103	Charlie	3	2	IT
103	Charlie	3	4	Marketing
104	David	NULL	1	HR
104	David	NULL	2	IT
104	David	NULL	4	Marketing

2. THETA JOIN (Using ON condition)

This is usually a condition-based join. We'll use Employees.DeptID = Departments.DeptID.

SELECT * FROM Employees E JOIN Departments D ON E.DeptID =
D.DeptID;

Output:

EmpID	EmpNam e	DeptID	DeptID	DeptNam e
101	Alice	1	1	HR
102	Bob	2	2	IT

Note: Only rows with matching DeptID are returned.

A **natural join** automatically joins using columns with the same name (DeptID in this case).

SELECT * FROM Employees NATURAL JOIN Departments;

Output:

EmpID	EmpNam	DeptID	DeptNam
	е		е

101 Alice 1	HR
-------------	----

Note: Same result as theta join here because both used DeptID as the common column.

- 1. LEFT OUTER JOIN
- 2. RIGHT OUTER JOIN
- 3. FULL OUTER JOIN

We'll use the same Employees and Departments tables.

(1) 1. LEFT OUTER JOIN

SELECT * FROM Employees E LEFT JOIN Departments D ON E.DeptID =
D.DeptID;

What it does:

- Returns all rows from Employees
- Adds matching Departments info when available
- If no match, fills with **NULLs**

Output:

EmpID	EmpNam e	DeptID	DeptID	DeptNam e
101	Alice	1	1	HR
102	Bob	2	2	IT
103	Charlie	3	NULL	NULL
104	David	NULL	NULL	NULL

Explanation:

- ullet Charlie's DeptID = 3 doesn't exist in Departments o NULLs
- $\bullet \quad \text{David has NULL DeptID} \rightarrow \text{can't match anything} \rightarrow \text{NULLs}$

2. RIGHT OUTER JOIN

Query:

sql

CopyEdit

SELECT *

FROM Employees E

RIGHT JOIN Departments D

ON E.DeptID = D.DeptID;

What it does:

- Returns **all rows** from Departments
- Adds matching Employees info when available
- If no match, fills with **NULLs**

Output:

EmpID	EmpNam e	DeptID	DeptID	DeptNam e
101	Alice	1	1	HR
102	Bob	2	2	IT
NULL	NULL	NULL	4	Marketing

Explanation:

ullet DeptID 4 (Marketing) has no employee o shows up with NULL employee data.

3. FULL OUTER JOIN

SELECT * FROM Employees E FULL OUTER JOIN Departments D
ON E.DeptID = D.DeptID;

What it does:

- Returns all rows from both tables
- Matches where it can
- Fills in NULLs where there's no match

Output:

EmpID	EmpNam e	DeptID	DeptID	DeptNam e
101	Alice	1	1	HR
102	Bob	2	2	IT
103	Charlie	3	NULL	NULL
104	David	NULL	NULL	NULL
NULL	NULL	NULL	4	Marketing

Explanation:

Combines left + right outer join
Everyone is included from both sides
UP TO HERE ALL JOINS USE ONLY TWO TABLES
A subquery (or nested query) is a query inside another query.
It usually helps answer questions like:
"Get all employees who earn more than the average salary"
"Get departments that don't have any employees"
"Find the employee with the highest salary"
Subqueries can be used in:
• SELECT
• FROM
• WHERE
• HAVING

Type Description

Scalar Subquery Returns a single value (1 row, 1 column)

Row Subquery Returns one row of multiple columns

Table Subquery Returns a full table (can be used like a normal

table)

Correlated Subquery

Depends on outer query for each row

★ Examples with Sample Table: Employees

EmpID	EmpName	DeptID	Salary
101	Alice	1	50000
102	Bob	2	60000
103	Charlie	1	70000
104	David	3	45000

Question: Get employees who earn more than the average salary.

```
SELECT * FROM Employees WHERE Salary > (
    SELECT AVG(Salary) FROM Employees );
```

- The inner query gets average salary: 56250
- Outer query returns Alice, Charlie, and Bob
- Question: Show each employee with the overall average salary.

```
SELECT EmpName, Salary, (SELECT AVG(Salary) FROM Employees) AS AvgSalary FROM Employees;
```

Question: Show employees from a temporary "high earners" table.

```
SELECT * FROM ( SELECT * FROM Employee WHERE Salary > 55000) AS HighEarners;
```

Explanation:

- Inner query creates a temporary table of high earners
- Outer query just selects from it

4. Correlated Subquery

Question: Get employees who earn more than anyone else in their department.

```
SELECT * FROM Employees E1 WHERE Salary > (SELECT MAX(Salary)
FROM Employees E2 WHERE E1.DeptID = E2.DeptID AND E1.EmpID <>
E2.EmpID );
```

Explanation:

- Inner query changes for each row
- Filters employees who have the highest salary in their own department

Practice Questions:

- 1. Get employees who earn more than the average salary.
- 2. List departments that have no employees (use subquery with NOT IN).
- 3. Show names of employees who earn less than Bob.
- 4. Find employee(s) with the highest salary.
- 5. Get employees whose salary is equal to the maximum in their department.

Sample Data

1. Employees Table

employee_id	employee_name	department_id
1	Alice	101
2	Bob	102

2. Departments Table

department_id department_name

101	Engineering
102	Marketing

3. ProjectAssignments Table

assignment_id	employee_id	project_id	assignment_start_date
1	1	201	2023-01-15
2	2	202	2023-02-01

4. Projects Table

project_id	project_name
201	Project X
202	Project Y

5. Salaries Table

salary_id employee_id	salary_amount
-----------------------	---------------

1	1	75000
2	2	65000

SQL Query

```
sql
```

SELECT

```
e.employee_name,
```

d.department_name,

p.project_name,

pa.assignment_start_date,

s.salary_amount

FROM Employees e

INNER JOIN Departments d ON e.department_id = d.department_id

INNER JOIN ProjectAssignments pa ON e.employee_id =

pa.employee_id

INNER JOIN Projects p ON pa.project_id = p.project_id

INNER JOIN Salaries s ON e.employee_id = s.employee_id;

Output

 $employee_name \quad department_name \quad project_name \quad assignment_start_date \quad salary_amount$

Alice Engineering Project X 2023-01-15 75000
--

Bob	Marketing	Project Y	2023-02-01	65000

You can copy and paste the above content directly into Google Docs, and it should maintain its formatting well!

1. Goal of the Query

This query retrieves a list of employees, their department, the projects they're assigned to, and their salaries. It combines data from five tables to create a comprehensive report.

2. Breakdown of Each Command

- a. FROM Employees e
 - Purpose: Start with the Employees table (aliased as e) as the primary source of data.
 - Why: The query focuses on employees, so this table acts as the "anchor" for joining other related tables.
- b. INNER JOIN Departments d ON e.department_id = d.department_id
 - Purpose: Link employees to their departments.
 - Why:
 - Every employee belongs to a department (assumed by the INNER JOIN).
 - If an employee has no department (department_id is NULL), they are excluded from results.
 - Columns like department_name come from this join.
- C. INNER JOIN ProjectAssignments pa ON e.employee_id = pa.employee_id
 - Purpose: Connect employees to their project assignments.
 - Why:
 - Each employee may be assigned to one or more projects (via the ProjectAssignments table).
 - INNER JOIN ensures only employees with project assignments are included.

If an employee has no projects, they are excluded from results.

- d. INNER JOIN Projects p ON pa.project_id = p.project_id
 - Purpose: Fetch details about the projects employees are assigned to.
 - Why:
 - The ProjectAssignments table only contains project_id, so we need the Projects table to get meaningful names (e.g., project_name).
 - INNER JOIN ensures only valid projects (with matching project_id) are included.
- e. INNER JOIN Salaries s ON e.employee_id = s.employee_id
 - Purpose: Retrieve salary information for each employee.
 - Why:
 - Assumes every employee has a salary record.
 - If an employee has no salary entry (e.g., Salaries table lacks their employee_id), they are excluded from results.

3. Why INNER JOIN Is Used Everywhere

Mandatory Relationships:

This guery assumes strict business rules:

- All employees must belong to a department.
- All employees must have at least one project assignment.
- All employees must have a salary record.
- INNER JOIN enforces these rules by excluding rows with missing relationships.

SQL GROUP BY and HAVING Guide

1. GROUP BY Clause

Purpose

Groups rows with the same values in specified columns into summary rows.

Syntax

```
sql
SELECT column1, aggregate_function(column2)
FROM table
GROUP BY column1;
```

2. HAVING Clause

Purpose

Filters groups created by GROUP BY based on a condition (like WHERE but for groups).

Key Difference from WHERE

- WHERE filters rows before grouping.
- HAVING filters groups after aggregation.

Syntax

```
sql
SELECT column1, aggregate_function(column2)
FROM table
GROUP BY column1
HAVING condition;
```

3. Aggregate Functions

Function Description Example

COUNT()	Counts the number of rows.	COUNT (order_id)
SUM()	Sums values in a numeric column.	SUM (amount)
AVG()	Calculates the average.	AVG(salary)
MIN()	Returns the smallest value.	MIN(price)
MAX()	Returns the largest value.	MAX (quantity)
STDDEV()	Standard deviation of values.	STDDEV(sales)
VARIANCE()	Variance of values.	VARIANCE (revenue)

4. Examples with Data, Queries, and Outputs

Sample Data: sales Table

sale_id	product	amount	region
1	Laptop	1000	East

2	Phone	500	East
3	Laptop	1200	West
4	Monitor	300	West
5	Phone	450	East

Example 1: COUNT with GROUP BY and HAVING

Goal: Find products sold more than once. Query:

guery

SELECT product, COUNT(*) AS total_sales
FROM sales

GROUP BY product
HAVING COUNT(*) > 1;

Output:

product total_sales

Laptop	2
Phone	2

Explanation:

- Groups data by product.
- COUNT (*) calculates total sales per product.
- HAVING filters out products with ≤ 1 sale.

Example 2: SUM with GROUP BY and HAVING

Goal: Find regions with total sales over \$1000.

Query:
sql

SELECT region, SUM(amount) AS total_revenue
FROM sales
GROUP BY region
HAVING SUM(amount) > 1000;

Output:

region total_revenue

East	1950
West	1500

Explanation:

- Groups data by region.
- SUM (amount) calculates total revenue per region.
- HAVING filters regions with total revenue ≤ \$1000.

Example 3: AVG with GROUP BY and HAVING

Goal: Find products with an average sale price ≥ \$700. Query:

sql

SELECT product, AVG(amount) AS avg_price FROM sales GROUP BY product

```
HAVING AVG(amount) >= 700;
Output:
               product
                                                      avg_price
                                     1100
 Laptop
Explanation:
   • Groups data by product.
   • AVG (amount) calculates average price per product.

    HAVING filters products with average price < $700.</li>

Example 4: MIN/MAX with GROUP BY and HAVING
Goal: Find regions where the smallest sale is < $400.
Query:
sql
SELECT region, MIN(amount) AS min_sale
FROM sales
GROUP BY region
HAVING MIN(amount) < 400;</pre>
Output:
                region
                                                      min_sale
```

300

Explanation:

West

- Groups data by region.
- MIN (amount) finds the smallest sale in each region.
- HAVING filters regions where the smallest sale ≥ \$400.

Example 5: Combining Multiple Aggregates

Output:

product	total_sales	max_sale
product	total_sales	IIIax_sait

Laptop	2200	1200

Explanation:

- Groups data by product.
- SUM (amount) and MAX (amount) calculate total and max sales.
- HAVING filters products that meet both conditions.

5. Key Rules

- 1. Columns in SELECT:
 - Must be in the GROUP BY clause or used in an aggregate function.
 - Example: SELECT region, SUM(amount) is valid if region is in GROUP BY.
- 2. Order of Execution:
 - ullet WHERE o GROUP BY o HAVING o ORDER BY.
- 3. Use Cases:
 - GROUP BY: Summarize data (e.g., sales by region).
 - HAVING: Filter aggregated results (e.g., total sales > \$1000).

6. HAVING vs WHFRF

HAVING Clause WHERE Clause

Filters groups after aggregation.	Filters rows before grouping.	
Used with aggregate functions.	Cannot use aggregate functions.	
Executed after GROUP BY.	Executed before GROUP BY.	

Copy and paste this entire structure into Google Docs. The markdown tables and code blocks will retain their formatting. Let me know if you need further adjustments!

SQL Guide: GROUP BY, Aggregates, and ORDER BY

Dataset: sales table

sale_id	product	amount	region	sale_date
1	Laptop	1000	East	2023-01-01
2	Phone	500	East	2023-01-02
3	Laptop	1200	West	2023-01-03

4	Monitor	300	West	2023-01-04
5	Phone	450	East	2023-01-05
6	Laptop	900	East	2023-01-06
7	Monitor	350	East	2023-01-07

Example 1: GROUP BY + SUM + ORDER BY

Goal

Total sales per product, ordered by highest sales first.

SQL Query

```
SELECT
  product,
  SUM(amount) AS total_sales
FROM sales
GROUP BY product
ORDER BY total_sales DESC;
```

Output

product total_sales

Laptop	3100
Phone	950
Monitor	650

Example 2: GROUP BY + AVG + HAVING + ORDER BY

Regions with average sale amount > \$400, ordered by region.

SQL Query

```
sql
SELECT
  region,
  AVG(amount) AS avg_sale
FROM sales
GROUP BY region
HAVING AVG(amount) > 400
ORDER BY region;
```

Output

region avg_sale

East 575.00	
-------------	--

West	750.00

1. Tables: Add gridlines for clarity.

Key Concepts

Concept	Description
---------	-------------

GROUP BY	Groups rows with identical values in specified columns.
Aggregates	Functions like SUM, AVG, COUNT to compute summary statistics.
HAVING	Filters groups after aggregation (used with GROUP BY).
ORDER BY	Sorts results in ascending (ASC) or descending (DESC) order.

LIMIT and OFFSET in MySQL

In MySQL, LIMIT and OFFSET are used to control the number of rows returned by a query. They are commonly used in pagination, where you fetch a subset of results instead of retrieving all data at once.

1. LIMIT

The LIMIT clause specifies the maximum number of rows to return.

```
Syntax:
sql
CopyEdit
SELECT column_names FROM table_name LIMIT number_of_rows;

Example:
sql
CopyEdit
SELECT * FROM customers LIMIT 5;
```

Explanation:

• This query returns only the first 5 rows from the customers table.

2. LIMIT with OFFSET

- The OFFSET clause skips a specified number of rows before starting to return rows.
- It is useful for pagination.

Syntax:

```
sql
CopyEdit
SELECT column_names FROM table_name LIMIT number_of_rows OFFSET
start_row;
SELECT column_names FROM table_name LIMIT start_row,
number_of_rows;
```

Note: The second syntax (LIMIT start, count) is a MySQL-specific way of using LIMIT and OFFSET.

Example:

sql

CopyEdit

```
SELECT * FROM customers LIMIT 5 OFFSET 10;
```

Explanation:

• This query skips the first 10 rows and then fetches the next 5 rows.

3. Pagination Example

Scenario: Fetch customer records in pages of 5 rows each

```
Page SQL Query

Page SELECT * FROM customers LIMIT 5
1 OFFSET 0;

Page SELECT * FROM customers LIMIT 5
2 OFFSET 5;

Page SELECT * FROM customers LIMIT 5
3 OFFSET 10;
```

Explanation:

- Page 1 fetches rows 1-5 (OFFSET 0, LIMIT 5)
- Page 2 fetches rows 6-10 (OFFSET 5, LIMIT 5)
- Page 3 fetches rows 11-15 (OFFSET 10, LIMIT 5)

4. Alternative: Using LIMIT with Start and Count

```
sql
```

CopyEdit

```
SELECT * FROM customers LIMIT 10, 5;
```

Explanation:

• This is the same as LIMIT 5 OFFSET 10, meaning it skips 10 rows and fetches 5 rows.

5. Sorting with LIMIT & OFFSET

Before paginating, always use ORDER BY to ensure a consistent order of results.

sql

CopyEdit

SELECT * FROM customers ORDER BY name ASC LIMIT 5 OFFSET 10;

Best Practice: Sorting avoids unexpected changes in data order when paginating.

6. Real-World Use Case: Fetching Products for an E-Commerce Website

Suppose you have a products table and want to display 10 products per page.

sql

CopyEdit

SELECT * FROM products ORDER BY product_name ASC LIMIT 10 OFFSET 20;

Explanation:

- This fetches Page 3 of results (assuming 10 results per page).
- OFFSET 20 means skipping first 20 products and fetching next 10.

Summary

Clause Purpose

LIMIT X Retrieves only X rows

LIMIT X OFFSET Y Skips Y rows and fetches X rows

LIMIT Y, X Another way to write LIMIT X OFFSET Y

(MySQL-specific)

Would you like an example with real sample data? 😊

Here's a structured exploration of LIMIT and OFFSET in combination with aggregate functions, GROUP BY, and ORDER BY using practical examples:

Sample Data: sales Table

sale_id	product	amount	region	customer_id
1	Laptop	1000	East	101
2	Phone	500	East	102
3	Laptop	1200	West	103
4	Monitor	300	West	104
5	Phone	450	East	101

6	Laptop	900	East	102
7	Monitor	350	East	103

Example 1: Top 3 Products by Total Sales

Goal: Show the top 3 products by total sales.

```
sql
SELECT product, SUM(amount) AS total_sales
FROM sales
GROUP BY product
ORDER BY total_sales DESC
LIMIT 3;
```

Output:

product total_sales

Laptop	3100
Phone	950
Monitor	650

Explanation:

- GROUP BY product groups sales by product.
- SUM(amount) calculates total sales per product.
- ORDER BY total_sales DESC sorts products from highest to lowest.
- LIMIT 3 returns only the top 3 rows.

Example 2: Second-Highest Average Sale by Region

Goal: Find the region with the second-highest average sale.

```
sql
SELECT region, AVG(amount) AS avg_sale
FROM sales
GROUP BY region
ORDER BY avg_sale DESC
LIMIT 1 OFFSET 1;
```

Output:

region	avg_sale	
East	575.00	

Explanation:

- GROUP BY region groups sales by region.
- AVG(amount) calculates average sales per region.
- ORDER BY avg_sale DESC sorts regions by highest average first.
- LIMIT 1 OFFSET 1 skips the first row (West: 750) and returns the second row (East: 575).

Example 3: Paginate Customers by Number of Orders

Goal: Show the 3rd and 4th customers by total orders (pagination).

```
sql
SELECT customer_id, COUNT(*) AS total_orders
FROM sales
GROUP BY customer_id
ORDER BY total_orders DESC
LIMIT 2 OFFSET 2;
```

Output:

customer_id total_orders

103	2
104	1

Explanation:

- GROUP BY customer_id groups sales by customer.
- COUNT(*) counts orders per customer.
- ORDER BY total_orders DESC sorts customers by most orders first.
- LIMIT 2 OFFSET 2 skips the first 2 rows and returns the next 2.

Example 4: Top 2 Regions by Total Sales (Excluding First)

Goal: Skip the top region and show the next 2 (if they exist).

```
sql
SELECT region, SUM(amount) AS total_sales
FROM sales
```

```
GROUP BY region
ORDER BY total_sales DESC
LIMIT 2 OFFSET 1;
```

Output:

	region		total_sales
East		2700	

Explanation:

- GROUP BY region groups sales by region.
- SUM(amount) calculates total sales per region.
- ORDER BY total_sales DESC sorts regions by highest sales first.
- LIMIT 2 OFFSET 1 skips the first region (West: 1500) and returns the next region (East: 2700).

(Only one region remains after OFFSET 1.)

Example 5: Products with the Lowest Single Sale

Goal: Find products with the smallest single sale, returning only 1 result.

```
sql
SELECT product, MIN(amount) AS min_sale
FROM sales
GROUP BY product
ORDER BY min_sale ASC
LIMIT 1;
```

Output:

I	product		min_sale
Monitor		300	

Explanation:

- GROUP BY product groups sales by product.
- MIN(amount) finds the smallest sale per product.
- ORDER BY min_sale ASC sorts products from smallest to largest minimum sale.
- LIMIT 1 returns the product with the smallest sale.

Key Takeaways

- LIMIT: Restricts the number of rows returned. Often used with ORDER BY to get "Top N" or "Bottom N" results.
- OFFSET: Skips a specified number of rows. Useful for pagination (e.g., page 2 of results).
- Order of Execution: GROUP BY \rightarrow Aggregations \rightarrow ORDER BY \rightarrow LIMIT/OFFSET.
- Common Use Cases: Paginating aggregated results, finding ranked results.

Pitfalls to Avoid

- Missing ORDER BY: Without sorting, LIMIT/OFFSET may return random rows.
- Large OFFSET: Skipping many rows can be slow; use keyset pagination for large datasets.
- Ambiguous Groups: Ensure GROUP BY includes all non-aggregated columns in SELECT.

Here's a structured guide on using the CASE statement in MySQL, complete with sample data, SQL queries, outputs, and explanations:

Practical Examples of Using the CASE Statement in MySQL

Sample Data: students Table

student_id	name	marks	gender
1	Alice	85	F
2	Bob	45	М
3	Charlie	92	М
4	Diana	72	F
5	Eva	60	F

Example 1: Categorize Grades

Goal: Assign grades based on marks.

SQL Query

```
sql
SELECT
name,
marks,
```

```
CASE

WHEN marks >= 90 THEN 'A'

WHEN marks >= 80 THEN 'B'

WHEN marks >= 60 THEN 'C'

ELSE 'Fail'

END AS grade

FROM students;
```

Output:

name marks grade

Alice	85	В
Bob	45	Fail
Charlie	92	A
Diana	72	С
Eva	60	С

Explanation:

The CASE statement evaluates marks and assigns grades conditionally.

Example 2: Conditional Aggregation

Goal: Count students by gender and pass/fail status.

SQL Query

```
sql
SELECT
  gender,
  COUNT(CASE WHEN marks >= 60 THEN 1 ELSE NULL END) AS passed,
  COUNT(CASE WHEN marks < 60 THEN 1 ELSE NULL END) AS failed
FROM students
GROUP BY gender;</pre>
```

Output:

gender	passed	failed
F	3	0
М	1	1

Explanation:

COUNT ignores NULL, so only valid conditions are counted.

- Female students (F): 3 passed, 0 failed.
- Male students (M): 1 passed (Charlie), 1 failed (Bob).

Example 3: Update Records Conditionally

Goal: Increase marks by 5 for students scoring below 50.

SQL Query

sql

```
UPDATE students
SET marks = CASE
  WHEN marks < 50 THEN marks + 5
  ELSE marks
END;</pre>
```

Result:

- Bob's marks become 50 (from 45).
- Other students' marks remain unchanged.

Example 4: Order by Custom Priority

Goal: Sort students with marks > 80 first, then others.

SQL Query

```
sql
SELECT name, marks
FROM students
ORDER BY
CASE
WHEN marks > 80 THEN 1
ELSE 2
END;
```

Output:

name marks

Alice	85

Charlie	92
Bob	45
Diana	72
Eva	60

Explanation:

Students with marks > 80 (Alice, Charlie) appear first.

Example 5: Pivot Data with CASE

Goal: Show total marks by gender and category (High/Low).

SQL Query

```
sql
SELECT
  gender,
  SUM(CASE WHEN marks >= 70 THEN marks ELSE 0 END) AS high_marks,
  SUM(CASE WHEN marks < 70 THEN marks ELSE 0 END) AS low_marks
FROM students
GROUP BY gender;</pre>
```

Output:

gender high_marks low_marks

F	157	60
М	92	45

Explanation:

- Female students: Alice (85) + Diana (72) = 157; Eva (60) = 60.
- Male students: Charlie (92); Bob (45) = 45.

Key Notes

Syntax:

```
sql
CASE
WHEN condition1 THEN result1
WHEN condition2 THEN result2
...
ELSE default_result
END
```

Use Cases:

- Categorizing data.
- Conditional updates.
- Custom sorting.
- Pivoting rows to columns.

ELSE Clause:

Optional but recommended to handle unmatched cases.

Feel free to copy and paste this content into Google Docs! Let me know if you'd like more advanced examples or any other modifications!

SQL Command to Determine Triangle Type

Sample Data: triangles Table

3	4	5
5	5	5
5	5	7
2	3	4
1	2	3

SQL Command

```
sql
SELECT
a, b, c,
-- Outer CASE: Check if it's a valid triangle
CASE
WHEN (a + b > c) AND (a + c > b) AND (b + c > a) THEN
-- Inner CASE: Determine the type of triangle
CASE
WHEN a = b AND b = c THEN 'Equilateral'
WHEN a = b OR b = c OR a = c THEN 'Isosceles'
WHEN POWER(a, 2) + POWER(b, 2) = POWER(c, 2) OR
```

Output

a b c triangle_type

3	4	5	Right-Angled
5	5	5	Equilateral
5	5	7	Isosceles
2	3	4	Scalene
1	2	3	Not a Triangle

Explanation

Outer CASE:

• Checks if the sides satisfy the triangle inequality theorem:

- a+b>c
- a+b>c
- a+c>b
- a+c>b
- b+c>a
- b+c>a
- If not satisfied, returns Not a Triangle.

Inner CASE:

- Equilateral: All sides are equal
- (a=b=c)
- (a=b=c).
- Isosceles: Exactly two sides are equal
- (a=b
- \bullet (a=b,
- b=c
- b=c, or
- a=c)
- \bullet a=c).
- Right-Angled: Satisfies the Pythagorean theorem (e.g.,
- a2+b2=c2
- Scalene: No sides are equal and not right-angled.

Key Notes

- 1. Order of Conditions: The CASE statement evaluates conditions sequentially. For example, Equilateral is checked before Isosceles because all equilateral triangles are technically isosceles, but we want the more specific classification first.
- 2. Right-Angled Check: The Pythagorean theorem is checked for all permutations of sides (since any side could be the hypotenuse).
- 3. Scalene: Falls under "none of the above" after checking other types.

Feel free to copy and paste this content into Google Docs! Let me know if you'd like to test with more examples or need any modifications!