

What is a Set?

A **Set** is a collection that does not allow duplicate elements. It is an unordered collection, meaning the elements are not stored in any particular order. Sets are commonly used when you need to ensure uniqueness of elements and perform operations like union, intersection, and difference.

In Java, the Set interface is part of the `java.util` package, and one of its most commonly used implementations is the **HashSet**.

HashSet in Java

- **HashSet** is a class that implements the Set interface.
 - It uses a **hash table** for storage, which provides constant-time performance for basic operations like add, remove, and contains.
 - **Key Features:**
 - Does not allow duplicate elements.
 - Allows null values.
 - Does not maintain insertion order.
 - Not thread-safe (for thread-safe operations, use `Collections.synchronizedSet` or `ConcurrentHashMap`).
-

Functions of HashSet

Here are the most commonly used methods of the HashSet class:

Method	Description
<code>boolean add(E e)</code>	Adds the specified element to the set if it is not already present.
<code>boolean remove(Object o)</code>	Removes the specified element from the set if it is present.
<code>void clear()</code>	Removes all elements from the set.
<code>boolean contains(Object o)</code>	Returns true if the set contains the specified element.
<code>boolean isEmpty()</code>	Returns true if the set contains no elements.
<code>int size()</code>	Returns the number of elements in the set.

<code>Iterator<E> iterator()</code>	Returns an iterator over the elements in the set.
<code>Object[] toArray()</code>	Returns an array containing all the elements in the set.
<code><T> T[] toArray(T[] a)</code>	Returns an array containing all the elements in the set (type-safe).
<code>boolean addAll(Collection<? extends E> c)</code>	Adds all elements from the specified collection to the set.
<code>boolean removeAll(Collection<?> c)</code>	Removes all elements from the set that are also in the specified collection.
<code>boolean retainAll(Collection<?> c)</code>	Retains only the elements in the set that are also in the specified collection.
<code>boolean containsAll(Collection<?> c)</code>	Returns true if the set contains all elements of the specified collection.
<code>Splitter<E> splitter()</code>	Creates a Splitter over the elements in the set.

Alternatives to HashSet

- **TreeSet**: Maintains elements in sorted order (uses a Red-Black tree).
- **LinkedHashSet**: Maintains insertion order (uses a hash table with a linked list).
- **ConcurrentHashMap**: Thread-safe alternative for concurrent environments.

Program Example: HashSet in Java

```
java
```

```
import java.util.HashSet;
import java.util.Iterator;
import java.util.Spliterator;
```

```

public class HashSetFunctionsExample {
    public static void main(String[] args) {
        // Create a HashSet
        HashSet<String> set = new HashSet<>();

        // 1. add(E e)
        set.add("Apple");
        set.add("Banana");
        set.add("Cherry");
        set.add("Apple"); // Duplicate element (will not be added)
        System.out.println("After adding elements: " + set);

        // 2. contains(Object o)
        System.out.println("Contains 'Banana'? " + set.contains("Banana"));

        // 3. remove(Object o)
        set.remove("Cherry");
        System.out.println("After removing 'Cherry': " + set);

        // 4. isEmpty()
        System.out.println("Is the set empty? " + set.isEmpty());

        // 5. size()
        System.out.println("Size of the set: " + set.size());

        // 6. iterator()
        System.out.println("Iterating over the set:");
        Iterator<String> iterator = set.iterator();
        while (iterator.hasNext()) {
            System.out.println(iterator.next());
        }

        // 7. toArray()
        Object[] array = set.toArray();
        System.out.println("Array from set: ");
        for (Object element : array) {
            System.out.println(element);
        }

        // 8. toArray(T[] a)
        String[] stringArray = new String[set.size()];
        set.toArray(stringArray);
        System.out.println("Type-safe array from set: ");
        for (String element : stringArray) {

```

```

        System.out.println(element);
    }

    // 9. addAll(Collection<? extends E> c)
    HashSet<String> anotherSet = new HashSet<>();
    anotherSet.add("Orange");
    anotherSet.add("Mango");
    set.addAll(anotherSet);
    System.out.println("After adding another set: " + set);

    // 10. removeAll(Collection<?> c)
    set.removeAll(anotherSet);
    System.out.println("After removing another set: " + set);

    // 11. retainAll(Collection<?> c)
    HashSet<String> retainSet = new HashSet<>();
    retainSet.add("Apple");
    retainSet.add("Banana");

    set.retainAll(retainSet);
    System.out.println("After retaining specific elements: " + set);

    // 12. containsAll(Collection<?> c)
    System.out.println("Does the set contain all elements of retainSet? " +
set.containsAll(retainSet));

    // 13. clear()
    set.clear();
    System.out.println("After clearing the set: " + set);

    // 14. spliterator()
    set.add("Grapes");
    set.add("Pineapple");
    Spliterator<String> spliterator = set.spliterator();
    System.out.println("Spliterator traversal:");
    spliterator.forEachRemaining(System.out::println);
}
}

```

OUTPUT 👍

After adding elements: [Apple, Banana, Cherry]
 Contains 'Banana'? true
 After removing 'Cherry': [Apple, Banana]

Is the set empty? false
Size of the set: 2
Iterating over the set:
Apple
Banana
Array from set:
Apple
Banana
Type-safe array from set:
Apple
Banana
After adding another set: [Apple, Banana, Orange, Mango]
After removing another set: [Apple, Banana]
After retaining specific elements: [Apple, Banana]
Does the set contain all elements of retainSet? true
After clearing the set: []
Splititerator traversal:
Grapes
Pineapple

Explanation of Each Function

1. **add(E e):**
 - Adds elements to the HashSet.
 - Duplicate elements (e.g., "Apple") are ignored.
2. **contains(Object o):**
 - Checks if the set contains a specific element ("Banana").
3. **remove(Object o):**
 - Removes a specific element ("Cherry") from the set.
4. **isEmpty():**
 - Checks if the set is empty.
5. **size():**
 - Returns the number of elements in the set.
6. **iterator():**
 - Provides an iterator to traverse the elements of the set.
7. **toArray():**
 - Converts the set into an array of Object type.
8. **toArray(T[] a):**
 - Converts the set into a type-safe array (String[] in this case).
9. **addAll(Collection<? extends E> c):**

- Adds all elements from another collection (anotherSet) to the set.
 - 10. **removeAll(Collection<?> c):**
 - Removes all elements from the set that are also present in another collection (anotherSet).
 - 11. **retainAll(Collection<?> c):**
 - Retains only the elements in the set that are also present in another collection (retainSet).
 - 12. **containsAll(Collection<?> c):**
 - Checks if the set contains all elements of another collection (retainSet).
 - 13. **clear():**
 - Removes all elements from the set.
 - 14. **splititerator():**
 - Creates a Splititerator for traversing the set. It is used for parallel processing.
-

Key Points

- The program demonstrates all the commonly used methods of HashSet.
- HashSet ensures that no duplicate elements are stored.
- The order of elements is not guaranteed in a HashSet.
- The program also shows how to convert a HashSet to an array and use a Splititerator.

This program is a comprehensive example of how to use HashSet in Java. You can modify and extend it to suit your specific use case!

Underlying Data Structure of HashSet

The **HashSet** in Java is implemented using a **hash table** data structure. Specifically:

- It uses a **HashMap** internally to store its elements.
- Each element in the HashSet is stored as a **key** in the underlying HashMap, and the value associated with each key is a constant (PRESENT), which is a dummy object.

How It Works

1. When you add an element to a HashSet, it calculates the **hash code** of the element using the `hashCode()` method.
2. The hash code is used to determine the **bucket** (index) where the element will be stored in the hash table.

3. If two elements have the same hash code (a **collision**), they are stored in the same bucket as a **linked list** (or a balanced tree in Java 8+ for performance optimization).
 4. To check if an element exists, the HashSet calculates the hash code and checks the corresponding bucket.
-

When to Use HashSet

Use HashSet when:

1. **Uniqueness is Required:**
 - You need to store a collection of unique elements (no duplicates).
2. **Fast Lookups are Needed:**
 - You need constant-time performance ($O(1)$) for add, remove, and contains operations.
3. **Order is Not Important:**
 - The order of elements does not matter.
4. **Memory Efficiency:**
 - You want to avoid storing duplicate elements, which saves memory.

Real-Life Examples of HashSet

1. **Unique Usernames:**
 - Storing a list of unique usernames in a system.
 2. **Tracking Visited URLs:**
 - Keeping track of visited URLs in a web crawler to avoid revisiting the same URL.
 3. **Removing Duplicates from a List:**
 - Removing duplicate entries from a list of email addresses.
 4. **Tagging System:**
 - Storing unique tags for blog posts or products.
-

Difference Between HashSet and ArrayList

Feature	HashSet	ArrayList
Underlying Structure	Hash table	Dynamic array
Duplicates	Does not allow duplicates	Allows duplicates

Order	No guaranteed order	Maintains insertion order
Performance	$O(1)$ for add, remove, and contains	$O(1)$ for add (amortized), $O(n)$ for search/remove
Use Case	Storing unique elements	Storing elements with order and duplicates
Memory Usage	More memory-efficient (no duplicates)	Less memory-efficient (allows duplicates)

Real-Life Example: HashSet vs ArrayList

Scenario: Storing Student IDs

- **Using HashSet:**
 - You want to store unique student IDs to ensure no duplicates.
 - Example:
 -

```
HashSet<String> studentIds = new HashSet<>();
studentIds.add("S101");
studentIds.add("S102");
studentIds.add("S101"); // Duplicate, will not be added
System.out.println(studentIds); // Output: [S101, S102]
```

Using ArrayList:

- You want to store student IDs in the order they were added, even if there are duplicates.
- Example:

```
ArrayList<String> studentIds = new ArrayList<>();
studentIds.add("S101");
studentIds.add("S102");
studentIds.add("S101"); // Duplicate, will be added
System.out.println(studentIds); // Output: [S101, S102, S101]
```

When to Use HashSet vs ArrayList

Use Case	HashSet	ArrayList
-----------------	----------------	------------------

Unique Elements	Use HashSet to ensure no duplicates.	Use ArrayList if duplicates are allowed.
Order Matters	Use LinkedHashSet if order matters.	Use ArrayList to maintain insertion order.
Fast Lookups	Use HashSet for fast lookups ($O(1)$).	ArrayList has slower lookups ($O(n)$).
Frequent Additions/Removals	Use HashSet for fast additions/removals.	ArrayList is slower for removals ($O(n)$).
Memory Efficiency	Use HashSet to avoid storing duplicates.	ArrayList may use more memory due to duplicates.

Real-Life Example: HashSet in a Social Media Application

Problem:

- A social media platform needs to store the unique usernames of all registered users.
- It should quickly check if a username is already taken during registration.

Solution:

- Use a HashSet to store usernames.
- Example:

```
HashSet<String> usernames = new HashSet<>();

// Adding usernames
usernames.add("john_doe");
usernames.add("jane_smith");
usernames.add("john_doe"); // Duplicate, will not be added

// Checking if a username exists
String newUsername = "john_doe";
if (usernames.contains(newUsername)) {
    System.out.println("Username already taken!");
} else {
    usernames.add(newUsername);
    System.out.println("Username registered successfully!");
}
```

```
}
```

Why HashSet?

- Ensures uniqueness of usernames.
 - Provides fast lookups to check if a username exists.
-

Real-Life Example: ArrayList in a Shopping Cart

Problem:

- An e-commerce platform needs to store items in a user's shopping cart.
- The order of items matters, and duplicates are allowed (e.g., multiple quantities of the same item).

Solution:

- Use an ArrayList to store items.
- Example:

```
ArrayList<String> cart = new ArrayList<>();
```

```
// Adding items to the cart
cart.add("Laptop");
cart.add("Mouse");
cart.add("Laptop"); // Duplicate, allowed
```

```
// Displaying the cart
System.out.println("Shopping Cart: " + cart);
```

Why ArrayList?

- Maintains the order of items.
 - Allows duplicates (e.g., multiple quantities of the same item).
-

Summary

- **HashSet:**
 - Underlying data structure: **Hash table**.
 - Use when you need **unique elements** and **fast lookups**.
 - Real-life example: Storing unique usernames.

- **ArrayList:**
 - Underlying data structure: **Dynamic array**.
 - Use when you need **ordered elements** and **allow duplicates**.
 - Real-life example: Storing items in a shopping cart.

Choose the appropriate data structure based on your requirements for **uniqueness**, **order**, and **performance**.

Here is a table summarizing the common functions available in **HashSet**, **TreeSet**, **LinkedHashSet**, and **EnumSet** in Java. The table indicates whether each function exists in the respective set type with "Yes" or "No".

Function	HashSet	TreeSet	LinkedHashSet	EnumSet
<code>add(E e)</code>	Yes	Yes	Yes	Yes
<code>addAll(Collection<? extends E> c)</code>	Yes	Yes	Yes	Yes
<code>clear()</code>	Yes	Yes	Yes	Yes
<code>clone()</code>	Yes	Yes	Yes	No
<code>contains(Object o)</code>	Yes	Yes	Yes	Yes
<code>containsAll(Collection<?> c)</code>	Yes	Yes	Yes	Yes
<code>isEmpty()</code>	Yes	Yes	Yes	Yes
<code>iterator()</code>	Yes	Yes	Yes	Yes
<code>remove(Object o)</code>	Yes	Yes	Yes	Yes
<code>removeAll(Collection<?> c)</code>	Yes	Yes	Yes	Yes
<code>retainAll(Collection<?> c)</code>	Yes	Yes	Yes	Yes

<code>size()</code>	Yes	Yes	Yes	Yes
<code>toArray()</code>	Yes	Yes	Yes	Yes
<code>toArray(T[] a)</code>	Yes	Yes	Yes	Yes
<code>comparator()</code>	No	Yes	No	No
<code>first()</code>	No	Yes	No	No
<code>last()</code>	No	Yes	No	No
<code>headSet(E toElement)</code>	No	Yes	No	No
<code>tailSet(E fromElement)</code>	No	Yes	No	No
<code>subSet(E fromElement, E toElement)</code>	No	Yes	No	No
<code>descendingIterator()</code>	No	Yes	No	No
<code>descendingSet()</code>	No	Yes	No	No
<code>pollFirst()</code>	No	Yes	No	No
<code>pollLast()</code>	No	Yes	No	No
<code>of()</code> (static factory methods)	No	No	No	Yes
<code>copyOf(Collection<E> c)</code>	No	No	No	Yes
<code>range(E from, E to)</code>	No	No	No	Yes

Notes:

1. **HashSet**: A general-purpose set implementation that uses a hash table for storage. It does not maintain any order.

2. **TreeSet**: A set implementation that uses a Red-Black tree for storage. It maintains elements in sorted order and provides additional navigation methods.
3. **LinkedHashSet**: A set implementation that maintains a linked list of the entries in the set, in the order in which they were inserted. It provides predictable iteration order.
4. **EnumSet**: A specialized set implementation for use with enum types. It is highly optimized for enums and provides additional static factory methods.

Let me know if you need further clarification!

LINKED HASHSET

LinkedHashSet Overview

A `LinkedHashSet` is a subclass of `HashSet` in Java that maintains a **linked list** of the elements in the set. This allows for **predictable iteration order**, which is the order in which elements were inserted into the set. Unlike `HashSet`, which does not guarantee any specific order, `LinkedHashSet` maintains **insertion order**.

Key Differences Between LinkedHashSet and HashSet

Feature	HashSet	LinkedHashSet
Ordering	No guaranteed order	Maintains insertion order
Performance	Slightly faster (no ordering overhead)	Slightly slower due to maintaining order
Internal Structure	Uses hash table	Uses hash table + doubly-linked list
Use Case	General-purpose set	When iteration order matters

Functions of LinkedHashSet

`LinkedHashSet` inherits all the methods from `HashSet` and adds the behavior of maintaining insertion order. Below is a list of its key methods:
When you want to avoid duplicates while preserving order

3. How Data is Stored

When you add an element to a `LinkedHashSet`, the following steps occur:

Step 1: Hash Calculation

- The `hashCode()` method of the element is called to compute its hash value.
- The hash value is used to determine the bucket (index) in the hash table where the element should be stored.

Step 2: Check for Duplicates

- If the bucket already contains elements, the `equals()` method is used to check if the element already exists in the set.
- If the element is already present, it is not added again (since sets do not allow duplicates).

Step 3: Add to Hash Table

- If the element is not already present, it is added to the appropriate bucket in the hash table.

Step 4: Add to Linked List

- A new node is created for the element.
- The node is added to the **end of the doubly-linked list** (to maintain insertion order).
- The `tail` reference is updated to point to this new node.

4. How Data is Accessed

When you iterate over a `LinkedHashSet`, the iteration follows the order of the **doubly-linked list**, not the hash table. This ensures that elements are returned in the order they were inserted.

Iteration Process

- Start from the head of the linked list.
 - Traverse the linked list using the `next` references of each node.
 - Stop when the `tail` is reached.
-

5. Internal Node Structure

Each node in the `LinkedHashSet` is represented by an instance of the `LinkedHashMap.Entry` class (since `LinkedHashSet` internally uses `LinkedHashMap`). The node contains:

- **Key:** The actual element stored in the set.
 - **Hash:** The hash value of the element.
 - **Next:** A reference to the next node in the hash table bucket (for collision resolution).
 - **Before:** A reference to the previous node in the linked list (for insertion order).
 - **After:** A reference to the next node in the linked list (for insertion order).
-

6. Example of Internal Storage

Let's say we add the following elements to a `LinkedHashSet`

```
LinkedHashSet<String> set = new LinkedHashSet<>();
set.add("Apple");
set.add("Banana");
set.add("Cherry");
```

Hash Table Structure

- The hash table might look like this (simplified for illustration):

Bucket Index	Elements
0	null
1	"Apple"
2	"Banana"
3	"Cherry"

Linked List Structure

- The doubly-linked list maintains the insertion order:
-

head -> "Apple" <-> "Banana" <-> "Cherry" <- tail

Iteration

- When you iterate over the set, the order will be: "Apple", "Banana", "Cherry".
-

7. Key Points About Internal Working

1. **Insertion Order:**
 2. **No Duplicates:**
 3. **Performance:**
 - **Add/Remove/Contains:** $O(1)$ on average (due to the hash table).
 - **Iteration:** $O(n)$ (due to the linked list, but faster than HashSet because it follows the linked list directly).
 4. **Memory Overhead:**
 - LinkedHashMap uses slightly more memory than HashSet because of the additional links (before and after) in each node.
-

8. Visual Representation

Hash Table + Linked List

Hash Table:

[0] -> null

[1] -> "Apple" -> null

[2] -> "Banana" -> null

[3] -> "Cherry" -> null

Linked List:

head -> "Apple" <-> "Banana" <-> "Cherry" <- tail

Iteration

- Iteration follows the linked list: "Apple" -> "Banana" -> "Cherry".
-

Summary

- LinkedHashMap uses a **hash table** for fast lookups and a **doubly-linked list** to maintain insertion order.
- It combines the benefits of HashSet (fast operations) with predictable iteration order.
- Internally, it relies on LinkedHashMap for storage and ordering.

- Use `LinkedHashSet` when you need a set with no duplicates and predictable iteration order.

Great question! Let me clarify the process of **how `LinkedHashSet` removes an element** in a more precise and step-by-step manner. The key point is that the **hash table and linked list are interconnected**, so there's no separate "searching" in the linked list. Here's how it works:

Detailed Removal Process

1. Hash Code Calculation

- When you call `remove(element)`, the `hashCode()` method of the element is called to compute its hash value.
 - This hash value is used to determine the **bucket (index)** in the hash table where the element might be stored.
-

2. Locate the Element in the Hash Table

- The system checks the corresponding bucket in the hash table.
 - If the bucket contains multiple elements (due to collisions), it uses the `equals()` method to find the exact match for the element.
-

3. Remove from the Hash Table

- Once the element is found in the hash table, it is removed from the bucket.
-

4. Remove from the Linked List

- Here's the important part: **the node in the hash table is the same node in the linked list**. There is no separate "searching" in the linked list because:
 - Each entry in the hash table is a **node** that contains:
 - The element (key).
 - References to the **previous** and **next** nodes in the linked list.
 - When the element is found in the hash table, you already have access to its linked list node.
- To remove the node from the linked list:

- Update the before and after references of the neighboring nodes to bypass the removed node.
 - If the removed node was the head or tail, update the head or tail reference accordingly.
-

Why No Separate Search in Linked List?

- The **hash table entry** and the **linked list node** are **the same object**. When you find the element in the hash table, you already have the linked list node.
 - This is why there's no need to "search" the linked list separately. The hash table lookup gives you direct access to the node in the linked list.
-

Example

Initial State

Hash Table:

[1] -> Node("Apple") -> null

[2] -> Node("Banana") -> null

[3] -> Node("Cherry") -> null

Linked List:

head -> Node("Apple") <-> Node("Banana") <-> Node("Cherry") <- tail

Remove "Banana"

1. Compute the hash code for "Banana" and locate bucket [2].
2. Find the node for "Banana" in the hash table.
3. Remove "Banana" from the hash table bucket.
4. Update the linked list:
 - Set Node("Apple").after = Node("Cherry").
 - Set Node("Cherry").before = Node("Apple").
5. Final state:

Hash Table:

[1] -> Node("Apple") -> null

[2] -> null

[3] -> Node("Cherry") -> null

Linked List:

```
head -> Node("Apple") <-> Node("Cherry") <- tail
```

Summary of Steps

1. Compute the hash code and locate the bucket in the hash table.
 2. Find the element in the bucket using `equals()`.
 3. Remove the element from the hash table.
 4. Remove the corresponding node from the linked list by updating its neighbors.
-

Key Takeaway

- The **hash table** and **linked list** are interconnected. The same node exists in both structures.
- When you find the element in the hash table, you already have access to its linked list node, so no separate search in the linked list is needed.
- This design ensures **efficient removal** with **O(1)** average time complexity.

ENUM SET

EnumSet Overview

EnumSet is a specialized set implementation in Java designed specifically for use with **enum types**. It is part of the `java.util` package and provides a highly efficient and compact way to work with sets of enum values. Unlike other set implementations like HashSet or TreeSet, EnumSet is optimized for enums and offers better performance.

Key Features of EnumSet

1. **Type-Safe**: Works only with a single enum type.
 2. **High Performance**: Uses bit vectors internally, making it extremely fast for enum operations.
 3. **Compact Storage**: Stores elements efficiently using bit flags.
 4. **Non-Thread Safe**: Like other sets, EnumSet is not thread-safe by default.
 5. **Nulls Not Allowed**: EnumSet does not allow null elements.
-

Advantages Over Other Sets

- **Faster:** Operations like add, remove, and contains are faster than HashSet or TreeSet.
 - **Memory Efficient:** Uses less memory compared to other sets.
 - **Specialized:** Designed specifically for enums, so it leverages the enum's ordinal values for storage.
-

Functions Specific to EnumSet

EnumSet inherits most of its methods from the Set interface, but it also provides some **factory methods** and **specialized behaviors** that are unique to it. Below are the functions specific to EnumSet:

2. Specialized Behaviors

- **Iteration Order:** Iterates over elements in the **natural order** of the enum (the order in which enum constants are declared).
- **Compact Representation:** Uses a **bit vector** internally, where each bit represents an enum constant. This makes operations like add, remove, and contains extremely fast.

```
import java.util.EnumSet;
```

```
public class EnumSetExample {  
    enum Day {  
        MONDAY, TUESDAY, WEDNESDAY, THURSDAY, FRIDAY, SATURDAY, SUNDAY  
    }  
}
```

```
    public static void main(String[] args) {  
        // Create an EnumSet containing all days  
        EnumSet<Day> allDays = EnumSet.allOf(Day.class);  
        System.out.println("All days: " + allDays);  
  
        // Create an EnumSet containing specific days  
        EnumSet<Day> weekdays = EnumSet.of(Day.MONDAY, Day.TUESDAY,  
        Day.WEDNESDAY, Day.THURSDAY, Day.FRIDAY);  
        System.out.println("Weekdays: " + weekdays);  
  
        // Create an EnumSet with a range of days  
        EnumSet<Day> weekend = EnumSet.range(Day.SATURDAY, Day.SUNDAY);  
        System.out.println("Weekend: " + weekend);  
    }  
}
```

```

        // Create a complement of the weekdays set
        EnumSet<Day> complement = EnumSet.complementOf weekdays);
        System.out.println("Complement of weekdays: " + complement);
    }
}

```

All days: [MONDAY, TUESDAY, WEDNESDAY, THURSDAY, FRIDAY, SATURDAY, SUNDAY]

Weekdays: [MONDAY, TUESDAY, WEDNESDAY, THURSDAY, FRIDAY]

Weekend: [SATURDAY, SUNDAY]

Complement of weekdays: [SATURDAY, SUNDAY]

Functions Not Available in Other Sets

1. **allOf(Class<E> elementType):**
 - Creates an EnumSet containing all elements of the specified enum type.
 - Not available in HashSet, TreeSet, or LinkedHashSet.
2. **noneOf(Class<E> elementType):**
 - Creates an empty EnumSet for the specified enum type.
 - Not available in other sets.
3. **range(E from, E to):**
 - Creates an EnumSet containing all elements in the range [from, to].
 - Not available in other sets.
4. **complementOf(EnumSet<E> s):**
 - Creates an EnumSet containing all elements not in the specified set.
 - Not available in other sets.
5. **of(E first, E... rest):**
 - Creates an EnumSet with the specified elements.
 - While other sets have of() methods (since Java 9), EnumSet's implementation is optimized for enums.

When to Use EnumSet

- When working with sets of enum values.
- When performance and memory efficiency are critical.
- When you need to perform operations like ranges or complements on enum sets.

Summary

- EnumSet is a specialized set for enum types.
- It provides **factory methods** like `allOf`, `noneOf`, `range`, and `complementOf` that are not available in other sets.
- It is **faster** and **more memory-efficient** than other sets for enum values.
- Use EnumSet when working with enums to leverage its optimizations.

TREE SET

TreeSet in Java

TreeSet is a class in Java that implements the NavigableSet interface, which is a subtype of the SortedSet interface. It stores elements in a **sorted order** (natural ordering or custom ordering via a Comparator). Internally, it uses a **Red-Black Tree** (a self-balancing binary search tree) to store elements, which ensures that operations like add, remove, and contains are performed in **$O(\log n)$** time.

Key Features of TreeSet

1. **Sorted Order:** Elements are stored in sorted order (ascending by default).
2. **No Duplicates:** Like all sets, TreeSet does not allow duplicate elements.
3. **Null Values:** TreeSet does not allow null values (throws `NullPointerException`).
4. **Thread Safety:** TreeSet is not thread-safe. For thread-safe operations, use `Collections.synchronizedSortedSet()`.

1. Constructors

Method	Description
<code>TreeSet()</code>	Constructs an empty TreeSet using the natural ordering of its elements.
<code>TreeSet(Comparator<? super E> comparator)</code>	Constructs an empty TreeSet ordered by the specified comparator.
<code>TreeSet(Collection<? extends E> c)</code>	Constructs a TreeSet containing the elements of the specified collection, ordered by natural ordering.

`TreeSet(SortedSet<E> s)` Constructs a `TreeSet` containing the same elements as the specified `SortedSet`, using the same ordering.

```
import java.util.*;
```

```
public class TreeSetExample {  
    public static void main(String[] args) {
```

```
        // 1. TreeSet() - Constructs an empty TreeSet using the natural ordering of its elements.
```

```
        TreeSet<Integer> treeSet1 = new TreeSet<>();  
        treeSet1.add(5);  
        treeSet1.add(2);  
        treeSet1.add(8);  
        System.out.println("TreeSet1 (Natural Ordering): " + treeSet1);
```

```
        // 2. TreeSet(Comparator<? super E> comparator) - Constructs an empty TreeSet ordered  
        by the specified comparator.
```

```
        TreeSet<Integer> treeSet2 = new TreeSet<>(Comparator.reverseOrder());  
        treeSet2.add(5);  
        treeSet2.add(2);  
        treeSet2.add(8);  
        System.out.println("TreeSet2 (Reverse Ordering): " + treeSet2);
```

```
        // 3. TreeSet(Collection<? extends E> c) - Constructs a TreeSet containing the elements of  
        the specified collection, ordered by natural ordering.
```

```
        List<Integer> list = Arrays.asList(7, 3, 9, 1);  
        TreeSet<Integer> treeSet3 = new TreeSet<>(list);  
        System.out.println("TreeSet3 (From Collection): " + treeSet3);
```

```
        // 4. TreeSet(SortedSet<E> s) - Constructs a TreeSet containing the same elements as the  
        specified SortedSet, using the same ordering.
```

```
        SortedSet<Integer> sortedSet = new TreeSet<>(Comparator.reverseOrder());  
        sortedSet.add(10);  
        sortedSet.add(4);  
        sortedSet.add(6);  
        TreeSet<Integer> treeSet4 = new TreeSet<>(sortedSet);  
        System.out.println("TreeSet4 (From SortedSet): " + treeSet4);  
    }  
}
```

TreeSet1 (Natural Ordering): [2, 5, 8]
TreeSet2 (Reverse Ordering): [8, 5, 2]
TreeSet3 (From Collection): [1, 3, 7, 9]
TreeSet4 (From SortedSet): [10, 6, 4]

3. Views

Method	Description
<code>Iterator<E> iterator()</code>	Returns an iterator over the elements in ascending order.
<code>Iterator<E> descendingIterator()</code>	Returns an iterator over the elements in descending order.
<code>NavigableSet<E> descendingSet()</code>	Returns a reverse-order view of the elements in the set.

```
import java.util.*;
```

```
public class TreeSetViewsExample {  
    public static void main(String[] args) {  
        // Create a TreeSet with some elements  
        TreeSet<Integer> treeSet = new TreeSet<>(Arrays.asList(10, 5, 15, 20, 25));  
  
        // 1. iterator() - Returns an iterator over the elements in ascending order  
        System.out.println("Ascending Order:");  
        Iterator<Integer> ascendingIterator = treeSet.iterator();  
        while (ascendingIterator.hasNext()) {  
            System.out.print(ascendingIterator.next() + " ");  
        }  
        System.out.println();  
  
        // 2. descendingIterator() - Returns an iterator over the elements in descending order  
        System.out.println("Descending Order:");  
        Iterator<Integer> descendingIterator = treeSet.descendingIterator();  
        while (descendingIterator.hasNext()) {  
            System.out.print(descendingIterator.next() + " ");  
        }  
        System.out.println();  
  
        // 3. descendingSet() - Returns a reverse-order view of the elements in the set  
        System.out.println("Reverse-Order View (DescendingSet:)");  
    }  
}
```



```

        NavigableSet<Integer> descendingSet = treeSet.descendingSet();
        for (Integer num : descendingSet) {
            System.out.print(num + " ");
        }
        System.out.println();
    }
}

```

Ascending Order:

5 10 15 20 25

Descending Order:

25 20 15 10 5

Reverse-Order View (DescendingSet):

25 20 15 10 5

. SortedSet Methods

Method	Description
Comparator<? super E> comparator()	Returns the comparator used to order the elements, or null if natural ordering is used.
E first()	Returns the first (lowest) element in the set.
E last()	Returns the last (highest) element in the set.
SortedSet<E> headSet(E toElement)	Returns a view of the portion of the set whose elements are strictly less than toElement.
SortedSet<E> tailSet(E fromElement)	Returns a view of the portion of the set whose elements are greater than or equal to fromElement.
SortedSet<E> subSet(E fromElement, E toElement)	Returns a view of the portion of the set whose elements range from fromElement (inclusive) to toElement (exclusive).

5. NavigableSet Methods

Method	Description
<code>E lower(E e)</code>	Returns the greatest element strictly less than the given element, or <code>null</code> if no such element exists.
<code>E floor(E e)</code>	Returns the greatest element less than or equal to the given element, or <code>null</code> if no such element exists.
<code>E ceiling(E e)</code>	Returns the smallest element greater than or equal to the given element, or <code>null</code> if no such element exists.
<code>E higher(E e)</code>	Returns the smallest element strictly greater than the given element, or <code>null</code> if no such element exists.
<code>E pollFirst()</code>	Retrieves and removes the first (lowest) element, or returns <code>null</code> if the set is empty.
<code>E pollLast()</code>	Retrieves and removes the last (highest) element, or returns <code>null</code> if the set is empty.
<code>NavigableSet<E> headSet(E toElement, boolean inclusive)</code>	Returns a view of the portion of the set whose elements are less than (or equal to, if <code>inclusive</code> is <code>true</code>) <code>toElement</code> .
<code>NavigableSet<E> tailSet(E fromElement, boolean inclusive)</code>	Returns a view of the portion of the set whose elements are greater than (or equal to, if <code>inclusive</code> is <code>true</code>) <code>fromElement</code> .
<code>NavigableSet<E> subSet(E fromElement, boolean fromInclusive, E toElement, boolean toInclusive)</code>	Returns a view of the portion of the set whose elements range from <code>fromElement</code> to <code>toElement</code> with optional inclusivity.

```
import java.util.*;
```

```
public class SortedSetMethodsExample {  
    public static void main(String[] args) {
```

```

// Create a TreeSet with some elements
TreeSet<Integer> treeSet = new TreeSet<>(Arrays.asList(10, 5, 15, 20, 25, 30, 35));

// 1. comparator() - Returns the comparator used to order the elements, or null if natural
ordering is used
Comparator<? super Integer> comparator = treeSet.comparator();
System.out.println("Comparator used: " + (comparator == null ? "Natural Ordering" :
comparator));

// 2. first() - Returns the first (lowest) element in the set
System.out.println("First element: " + treeSet.first());

// 3. last() - Returns the last (highest) element in the set
System.out.println("Last element: " + treeSet.last());

// 4. headSet(E toElement) - Returns a view of the portion of the set whose elements are
strictly less than toElement
SortedSet<Integer> headSet = treeSet.headSet(20);
System.out.println("HeadSet (elements < 20): " + headSet);

// 5. tailSet(E fromElement) - Returns a view of the portion of the set whose elements are
greater than or equal to fromElement
SortedSet<Integer> tailSet = treeSet.tailSet(20);
System.out.println("TailSet (elements >= 20): " + tailSet);

// 6. subSet(E fromElement, E toElement) - Returns a view of the portion of the set whose
elements range from fromElement (inclusive) to toElement (exclusive)
SortedSet<Integer> subSet = treeSet.subSet(15, 30);
System.out.println("SubSet (elements >= 15 and < 30): " + subSet);
}
}

```

```

Comparator used: Natural Ordering
First element: 5
Last element: 35
HeadSet (elements < 20): [5, 10, 15]
TailSet (elements >= 20): [20, 25, 30, 35]
SubSet (elements >= 15 and < 30): [15, 20, 25]

```

5. NavigableSet Methods

Method	Description
--------	-------------

<code>E lower(E e)</code>	Returns the greatest element strictly less than the given element, or <code>null</code> if no such element exists.
<code>E floor(E e)</code>	Returns the greatest element less than or equal to the given element, or <code>null</code> if no such element exists.
<code>E ceiling(E e)</code>	Returns the smallest element greater than or equal to the given element, or <code>null</code> if no such element exists.
<code>E higher(E e)</code>	Returns the smallest element strictly greater than the given element, or <code>null</code> if no such element exists.
<code>E pollFirst()</code>	Retrieves and removes the first (lowest) element, or returns <code>null</code> if the set is empty.
<code>E pollLast()</code>	Retrieves and removes the last (highest) element, or returns <code>null</code> if the set is empty.
<code>NavigableSet<E> headSet(E toElement, boolean inclusive)</code>	Returns a view of the portion of the set whose elements are less than (or equal to, if <code>inclusive</code> is <code>true</code>) <code>toElement</code> .
<code>NavigableSet<E> tailSet(E fromElement, boolean inclusive)</code>	Returns a view of the portion of the set whose elements are greater than (or equal to, if <code>inclusive</code> is <code>true</code>) <code>fromElement</code> .
<code>NavigableSet<E> subSet(E fromElement, boolean fromInclusive, E toElement, boolean toInclusive)</code>	Returns a view of the portion of the set whose elements range from <code>fromElement</code> to <code>toElement</code> with optional inclusivity.

```
import java.util.*;
```

```
public class NavigableSetMethodsExample {
    public static void main(String[] args) {
        // Create a TreeSet with some elements
        TreeSet<Integer> treeSet = new TreeSet<>(Arrays.asList(10, 5, 15, 20, 25, 30, 35));

        System.out.println("Original TreeSet: " + treeSet);

        // 1. lower(E e) - Returns the greatest element strictly less than the given element
```

```

System.out.println("Lower than 20: " + treeSet.lower(20)); // 15

// 2. floor(E e) - Returns the greatest element less than or equal to the given element
System.out.println("Floor of 20: " + treeSet.floor(20)); // 20

// 3. ceiling(E e) - Returns the smallest element greater than or equal to the given element
System.out.println("Ceiling of 22: " + treeSet.ceiling(22)); // 25

// 4. higher(E e) - Returns the smallest element strictly greater than the given element
System.out.println("Higher than 20: " + treeSet.higher(20)); // 25

// 5. pollFirst() - Retrieves and removes the first (lowest) element
System.out.println("PollFirst: " + treeSet.pollFirst()); // 5
System.out.println("TreeSet after pollFirst: " + treeSet);

// 6. pollLast() - Retrieves and removes the last (highest) element
System.out.println("PollLast: " + treeSet.pollLast()); // 35
System.out.println("TreeSet after pollLast: " + treeSet);

// 7. headSet(E toElement, boolean inclusive) - Returns a view of the portion of the set
whose elements are less than (or equal to, if inclusive is true) toElement
NavigableSet<Integer> headSet = treeSet.headSet(20, true);
System.out.println("HeadSet (elements <= 20): " + headSet);

// 8. tailSet(E fromElement, boolean inclusive) - Returns a view of the portion of the set
whose elements are greater than (or equal to, if inclusive is true) fromElement
NavigableSet<Integer> tailSet = treeSet.tailSet(20, true);
System.out.println("TailSet (elements >= 20): " + tailSet);

// 9. subSet(E fromElement, boolean fromInclusive, E toElement, boolean toInclusive) -
Returns a view of the portion of the set whose elements range from fromElement to toElement
with optional inclusivity
NavigableSet<Integer> subSet = treeSet.subSet(15, true, 25, false);
System.out.println("SubSet (elements >= 15 and < 25): " + subSet);
}
}

```

Original TreeSet: [5, 10, 15, 20, 25, 30, 35]
 Lower than 20: 15
 Floor of 20: 20
 Ceiling of 22: 25
 Higher than 20: 25
 PollFirst: 5
 TreeSet after pollFirst: [10, 15, 20, 25, 30, 35]

PollLast: 35

TreeSet after pollLast: [10, 15, 20, 25, 30]

HeadSet (elements ≤ 20): [10, 15, 20]

TailSet (elements ≥ 20): [20, 25, 30]

SubSet (elements ≥ 15 and < 25): [15, 20]

Skip List Set Overview

A **Skip List Set** is a probabilistic data structure that allows for fast search, insertion, and deletion operations. It is an alternative to balanced trees (like AVL or Red-Black trees) and provides **average-case $O(\log n)$** time complexity for these operations. The Skip List Set is implemented in Java as part of the `ConcurrentSkipListSet` class, which is a thread-safe, sorted set.

Key Features of Skip List Set

1. **Sorted Order:** Elements are stored in sorted order.
 2. **Probabilistic Structure:** Uses multiple layers of linked lists to skip over elements, reducing the number of comparisons needed.
 3. **Thread-Safe:** `ConcurrentSkipListSet` is thread-safe and supports concurrent access.
 4. **Efficient Operations:** Provides **$O(\log n)$** average-case time complexity for search, insert, and delete operations.
-

Internal Working of Skip List Set

1. Structure of a Skip List

A Skip List is composed of multiple layers of linked lists:

- **Base Layer:** A standard linked list containing all the elements in sorted order.
- **Upper Layers:** Sparse linked lists that "skip" over elements, allowing for faster traversal.

Each node in the Skip List contains:

- A **value** (the element stored in the set).
 - **Forward pointers:** References to the next nodes at each level.
 - A **height:** The number of levels the node is part of.
-

2. How It Works

- **Search:**
 - Start at the topmost layer of the Skip List.
 - Move forward as long as the next element is less than the target.
 - If the next element is greater, move down to the next layer and repeat.
 - Continue until the element is found or determined to be absent.
 - **Insertion:**
 - Use the search process to find the position where the new element should be inserted.
 - Randomly determine the height of the new node (number of layers it will be part of).
 - Insert the new node into all layers up to its height.
 - **Deletion:**
 - Use the search process to locate the node to be deleted.
 - Remove the node from all layers it is part of.
-

3. Probabilistic Nature

- The height of each node is determined randomly using a probabilistic algorithm (e.g., flipping a coin).
 - This randomness ensures that the Skip List remains balanced on average, providing **$O(\log n)$** performance.
-

(Due to technical issues, the search service is temporarily unavailable.)

When two threads attempt to perform operations (e.g., insert, delete, or update) on the **same element** in a `ConcurrentSkipListSet`, the behavior depends on the **thread-safety mechanisms** implemented in the `ConcurrentSkipListSet`. Here's a detailed explanation of what happens:

1. Thread-Safety in `ConcurrentSkipListSet`

`ConcurrentSkipListSet` is designed to handle concurrent access safely. It uses **lock-free algorithms** and **atomic operations** to ensure that multiple threads can perform operations concurrently without causing data corruption or inconsistencies.

2. Scenarios and Outcomes

Scenario 1: Two Threads Insert the Same Element

- If two threads try to **insert the same element** simultaneously:

- Only **one thread** will succeed in inserting the element.
- The other thread will detect that the element already exists (due to the nature of a set, which does not allow duplicates) and will **not insert it again**.
- The `add()` method will return `true` for the thread that successfully inserts the element and `false` for the other thread.

Scenario 2: One Thread Inserts and Another Thread Deletes the Same Element

- If one thread tries to **insert** an element while another thread tries to **delete** the same element:
 - The outcome depends on the **order of operations**:
 - If the **delete operation** happens first, the element will be removed, and the **insert operation** will succeed in adding the element.
 - If the **insert operation** happens first, the element will be added, and the **delete operation** will remove it.
 - The operations are **atomic**, so there will be no intermediate inconsistent state.

Scenario 3: Two Threads Update the Same Element

- Since `ConcurrentSkipListSet` is a **set**, it does not support direct updates to elements. Instead, you would need to **remove** the old element and **insert** a new one.
- If two threads try to **remove and insert** the same element simultaneously:
 - The operations will be performed atomically, and the final state will depend on the order of operations.
 - For example:
 - Thread 1 removes element A.
 - Thread 2 removes element A (but it's already removed).
 - Thread 1 inserts a new element A'.
 - Thread 2 inserts a new element A''.
 - The final set will contain either A' or A'', depending on which thread's insert operation succeeds last.

Scenario 4: Two Threads Perform a Read Operation

- If two threads perform a **read operation** (e.g., `contains()` or iteration) on the same element:
 - Both threads will see a **consistent view** of the set at the time of their operation.

- Read operations do not modify the set, so they can proceed concurrently without any issues.
-

3. How ConcurrentSkipListSet Ensures Thread-Safety

- **Lock-Free Algorithms:** ConcurrentSkipListSet uses **compare-and-swap (CAS)** operations to ensure atomicity without locks.
 - **Immutable Nodes:** Once a node is added to the Skip List, it is not modified. Updates are performed by creating new nodes and linking them atomically.
 - **Memory Consistency:** Changes made by one thread are guaranteed to be visible to other threads due to the **happens-before relationship** enforced by the ConcurrentSkipListSet.
-

4. Example: Concurrent Operations on ConcurrentSkipListSet

```
import java.util.concurrent.ConcurrentSkipListSet;

public class ConcurrentSkipListSetExample {
    public static void main(String[] args) throws InterruptedException {
        ConcurrentSkipListSet<Integer> skipListSet = new ConcurrentSkipListSet<>();

        // Thread 1: Insert elements
        Thread thread1 = new Thread(() -> {
            for (int i = 0; i < 10; i++) {
                skipListSet.add(i);
                System.out.println("Thread 1 added: " + i);
            }
        });

        // Thread 2: Insert elements (including duplicates)
        Thread thread2 = new Thread(() -> {
            for (int i = 5; i < 15; i++) {
                boolean added = skipListSet.add(i);
                System.out.println("Thread 2 added: " + i + " (success: " + added + ")");
            }
        });

        // Start both threads
        thread1.start();
        thread2.start();
    }
}
```

```
// Wait for threads to finish
thread1.join();
thread2.join();

// Display the final set
System.out.println("Final SkipListSet: " + skipListSet);
}
}
```

Thread 1 added: 0
Thread 1 added: 1
Thread 2 added: 5 (success: true)
Thread 1 added: 2
Thread 2 added: 6 (success: true)
Thread 1 added: 3
Thread 2 added: 7 (success: true)
Thread 1 added: 4
Thread 2 added: 8 (success: true)
Thread 1 added: 5 (success: false)
Thread 2 added: 9 (success: true)
Thread 1 added: 6 (success: false)
Thread 2 added: 10 (success: true)
Thread 1 added: 7 (success: false)
Thread 2 added: 11 (success: true)
Thread 1 added: 8 (success: false)
Thread 2 added: 12 (success: true)
Thread 1 added: 9 (success: false)
Thread 2 added: 13 (success: true)
Thread 2 added: 14 (success: true)
Final SkipListSet: [0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14]

Explanation of the Output

- Both threads attempt to insert elements into the `ConcurrentSkipListSet`.
- When both threads try to insert the **same element** (e.g., 5), only one thread succeeds, and the other thread's `add()` method returns `false`.
- The final set contains all unique elements from both threads, and the order is maintained.

Key Takeaways

- `ConcurrentSkipListSet` ensures **thread-safety** using lock-free algorithms and atomic operations.
- When two threads perform operations on the same element:
 - Only one thread will succeed in inserting or deleting the element.
 - The final state of the set will be consistent and correct.
- Read operations are always safe and do not interfere with each other.

This makes `ConcurrentSkipListSet` a reliable choice for **concurrent, sorted sets** in multi-threaded applications.

(Due to technical issues, the search service is temporarily unavailable.)

In an `ArrayList`, when you **update a number** (or any element), the **existing object is modified**, and **no new object is created**. Here's a simple explanation:

What Happens When You Update an Element in an `ArrayList`?

1. **Direct Update:**
 - The `ArrayList` stores references to objects (e.g., `Integer`, `String`, etc.).
 - When you update an element using `set(index, element)`, the reference at that index is **replaced** with the new element.
 - If the element is a primitive (e.g., `int`), it is autoboxed into its wrapper class (e.g., `Integer`), but the reference in the `ArrayList` is still updated directly.
 2. **No New Object for the `ArrayList`:**
 - The `ArrayList` itself is **not recreated**; only the specific element at the given index is updated.
-

(Due to technical issues, the search service is temporarily unavailable.)

Below is a comparison of the functions available in `ConcurrentSkipListSet` and `CopyOnWriteArraySet`. The functions are listed on the left, and Yes/No indicates whether the function is available in each set.

Functions in `ConcurrentSkipListSet` vs `CopyOnWriteArraySet`

Function	ConcurrentSkipListSet	CopyOnWriteArraySet
<code>add(E e)</code>	Yes	Yes
<code>addAll(Collection<? extends E> c)</code>	Yes	Yes
<code>clear()</code>	Yes	Yes
<code>contains(Object o)</code>	Yes	Yes
<code>containsAll(Collection<?> c)</code>	Yes	Yes
<code>equals(Object o)</code>	Yes	Yes
<code>hashCode()</code>	Yes	Yes
<code>isEmpty()</code>	Yes	Yes
<code>iterator()</code>	Yes	Yes
<code>parallelStream()</code>	Yes	Yes
<code>remove(Object o)</code>	Yes	Yes

<code>removeAll(Collection<?> c)</code>	Yes	Yes
---	-----	-----

<code>removeIf(Predicate<? super E> filter)</code>	Yes	Yes
--	-----	-----

<code>retainAll(Collection<?> c)</code>	Yes	Yes
---	-----	-----

<code>size()</code>	Yes	Yes
---------------------	-----	-----

<code>splititerator()</code>	Yes	Yes
------------------------------	-----	-----

<code>stream()</code>	Yes	Yes
-----------------------	-----	-----

<code>toArray()</code>	Yes	Yes
------------------------	-----	-----

<code>toArray(T[] a)</code>	Yes	Yes
-----------------------------	-----	-----

<code>toString()</code>	Yes	Yes
-------------------------	-----	-----

SortedSet Functions

<code>comparator()</code>	Yes	No
---------------------------	-----	----

<code>first()</code>	Yes	No
----------------------	-----	----

<code>last()</code>	Yes	No
---------------------	------------	-----------

<code>headSet(E toElement)</code>	Yes	No
-----------------------------------	------------	-----------

<code>tailSet(E fromElement)</code>	Yes	No
-------------------------------------	------------	-----------

<code>subSet(E fromElement, E toElement)</code>	Yes	No
---	------------	-----------

Concurrent Functions

<code>descendingIterator()</code>	Yes	No
-----------------------------------	------------	-----------

<code>descendingSet()</code>	Yes	No
------------------------------	------------	-----------

<code>ceiling(E e)</code>	Yes	No
---------------------------	------------	-----------

<code>floor(E e)</code>	Yes	No
-------------------------	------------	-----------

<code>higher(E e)</code>	Yes	No
--------------------------	------------	-----------

<code>lower(E e)</code>	Yes	No
-------------------------	------------	-----------

<code>pollFirst()</code>	Yes	No
--------------------------	------------	-----------

`pollLast()`

Yes

No

Key Differences

1. SortedSet Functions

- **ConcurrentSkipListSet implements the SortedSet interface, so it provides functions like:**
 - `comparator()`
 - `first()`
 - `last()`
 - `headSet()`
 - `tailSet()`
 - `subSet()`
- **CopyOnWriteArraySet does not implement SortedSet, so these functions are not available.**

2. Concurrent Functions

- **ConcurrentSkipListSet provides additional concurrent functions like:**
 - `descendingIterator()`
 - `descendingSet()`
 - `ceiling()`
 - `floor()`
 - `higher()`
 - `lower()`
 - `pollFirst()`
 - `pollLast()`
- **These functions are not available in CopyOnWriteArraySet.**

3. Performance Characteristics

- **ConcurrentSkipListSet:**
 - Provides sorted order and logarithmic-time operations.
 - Ideal for concurrent scenarios requiring sorted data.
 - **CopyOnWriteArraySet:**
 - Provides immutable snapshots of the data.
 - Ideal for scenarios with frequent reads and rare writes.
-

When to Use Which?

- **Use ConcurrentSkipListSet if:**
 - You need sorted data.
 - You require efficient concurrent operations.
 - You need additional functions like `ceiling()`, `floor()`, etc.
 - **Use CopyOnWriteArraySet if:**
 - You have frequent reads and rare writes.
 - You do not need sorted data.
 - You want simplicity and thread-safety without additional overhead.
-

Summary

- Both **ConcurrentSkipListSet** and **CopyOnWriteArraySet** provide thread-safe implementations of sets.
- **ConcurrentSkipListSet** offers sorted data and additional concurrent functions.
- **CopyOnWriteArraySet** is simpler and optimized for read-heavy workloads.

CopyOnWriteArraySet Overview

CopyOnWriteArraySet is a thread-safe implementation of the **Set** interface in Java. It is part of the `java.util.concurrent` package and is designed for scenarios where read operations significantly outnumber write operations. It achieves thread-safety by creating a new copy of the underlying array whenever the set is modified, ensuring that read operations are not blocked by write operations.

Key Features of CopyOnWriteArraySet

1. Thread-Safe:

- All operations are thread-safe without the need for explicit synchronization.
- Read operations are lock-free and always see a consistent snapshot of the data.

2. Immutability During Iteration:

- Iterators operate on a snapshot of the data at the time the iterator was created.
- This ensures that the iterator does not throw a `ConcurrentModificationException`.

3. Write Overhead:

- Write operations (e.g., add, remove) are expensive because they involve creating a new copy of the underlying array.
- This makes `CopyOnWriteArraySet` suitable for read-heavy workloads.

4. No Sorted Order:

- Elements are not stored in sorted order (unlike `ConcurrentSkipListSet`).

5. No Null Elements:

- `CopyOnWriteArraySet` does not allow null elements.
-

Internal Working of CopyOnWriteArraySet

1. Underlying Data Structure

- `CopyOnWriteArraySet` is implemented using a `CopyOnWriteArrayList` internally.
- The `CopyOnWriteArrayList` stores the elements of the set in an array.

2. Write Operations

- When a write operation (e.g., `add`, `remove`) is performed:
 1. A new copy of the underlying array is created.
 2. The modification (e.g., adding or removing an element) is applied to the new copy.
 3. The reference to the underlying array is updated to point to the new copy.

3. Read Operations

- Read operations (e.g., `contains`, `iterator`) operate on the current snapshot of the array.
- Since the array is immutable during iteration, read operations are fast and thread-safe.

4. Iterators

- Iterators operate on the snapshot of the array at the time the iterator was created.
- Changes to the set after the iterator is created are not visible to the iterator.

Example Usage

```
import java.util.concurrent.CopyOnWriteArraySet;
```

```
public class CopyOnWriteArraySetExample {
```

```
public static void main(String[] args) {  
  
    // Create a CopyOnWriteArraySet  
  
    CopyOnWriteArraySet<String> set = new CopyOnWriteArraySet<>();  
  
  
    // Add elements  
  
    set.add("Apple");  
  
    set.add("Banana");  
  
    set.add("Cherry");  
  
  
    // Display the set  
  
    System.out.println("Initial Set: " + set);  
  
  
    // Iterate over the set  
  
    System.out.println("Iterating over the set:");  
  
    for (String fruit : set) {  
  
        System.out.println(fruit);  
  
    }  
}
```

```
// Add an element while iterating
```

```
set.add("Date");
```

```
// Display the updated set
```

```
System.out.println("Updated Set: " + set);
```

```
}
```

```
}
```

Initial Set: [Apple, Banana, Cherry]

Iterating over the set:

Apple

Banana

Cherry

Updated Set: [Apple, Banana, Cherry, Date]

Key Points from the Example

- The initial set contains [Apple, Banana, Cherry].
- During iteration, the set is immutable, so the iterator does not see the newly added element "Date".
- After iteration, the set is updated to include "Date".

Advantages of `CopyOnWriteArraySet`

1. Thread-Safety:

- No need for explicit synchronization.
- Read operations are always consistent.

2. No `ConcurrentModificationException`:

- Iterators operate on a snapshot, so they never throw `ConcurrentModificationException`.

3. Simple and Predictable:

- Easy to use and understand.
-

Disadvantages of `CopyOnWriteArraySet`

1. Write Overhead:

- Write operations are expensive due to the creation of a new array copy.

2. Memory Usage:

- Frequent modifications can lead to high memory usage due to repeated array copying.

3. No Sorted Order:

- Elements are not stored in sorted order.
-

When to Use `CopyOnWriteArraySet`

- Read-Heavy Workloads:

- When the number of read operations far exceeds the number of write operations.
 - Event Listeners:
 - Commonly used for managing lists of event listeners in multi-threaded environments.
 - Immutable Snapshots:
 - When you need consistent snapshots of the data for iteration or analysis.
-

Comparison with `ConcurrentSkipListSet`

Feature	<code>CopyOnWriteArraySet</code>	<code>ConcurrentSkipListSet</code>
Thread-Safety	Yes	Yes
Write Performance	Slow (due to array copying)	Fast (lock-free updates)
Read Performance	Fast (lock-free)	Fast (lock-free)
Sorted Order	No	Yes
Iterator Behavior	Snapshot at creation time	Reflects current state

Use Case

Read-heavy workloads

Concurrent sorted
sets

Summary

- `CopyOnWriteArraySet` is a thread-safe set that creates a new copy of the underlying array on every modification.
- It is ideal for read-heavy workloads where writes are infrequent.
- It provides consistent snapshots for iteration and is easy to use in multi-threaded environments.
- However, it is not suitable for write-heavy workloads due to the overhead of array copying.